

UNIVERSITÀ DEGLI STUDI DI VERONA

DEPARTMENT OF ECONOMICS

Master's Degree in Economics and Data Analysis

**DO MACHINE LEARNING MODELS IMPROVE
STOCK RETURN PREDICTION?**

**Evidence from S&P 500 Constituent Markets and Dimension Sensitivity
(2015–2024)**

Supervisor:

Prof. Giuseppina Chesini

Candidate:

Murtuza Yusuf Rangwala

Matricola: VR508566

Academic Year 2024/2025

Abstract

Do complex machine learning algorithms actually outperform classical econometric benchmarks when predicting daily US stock returns? I address this empirical question by testing a broad range of quantitative models across a 10-year point-in-time panel of S&P 500 constituent equities from 2015 through 2024 (2,516 daily cross-sections). The predictor universe comprises 59 characteristics spanning technical indicators, quarterly Bloomberg accounting ratios, macroeconomic state variables, news/social sentiment scores, and rolling factor risk exposures.

To prevent lookahead leakage, model parameters and hyperparameter search spaces are optimized strictly inside a 5-fold expanding-window walk-forward validation protocol across the 2020–2024 out-of-sample period. Empirical results confirm that deep sequence models outperform linear estimators (Fama-MacBeth OLS, ElasticNet, LASSO) and static decision tree ensembles (Random Forest, XGBoost). My custom **Temporal Fusion Deep Multimodal Gated Attention (TFDMGA)** architecture delivers the highest predictive accuracy, recording an out-of-sample daily rank Information Coefficient (IC) of +0.0348, an Information Ratio ($ICIR$) of 3.12, a ROC AUC of 0.6120, and a directional accuracy of 56.84%. A Diebold-Mariano test with Harvey-Leybourne-Newbold small-sample corrections verifies that these forecast gains over baseline PyTorch LSTMs are statistically significant ($DM = 2.41, p = 0.016$).

When deployed in simulated daily trading portfolios incorporating 10 bps round-trip transaction frictions and a 1-day execution buffer, an initial \$1,000 USD deposit starting in January 2020 grows to \$3,120.99 USD under a baseline LSTM long strategy. Adding a 2:1 Take-Profit/Stop-Loss risk overlay expands capital accumulation to \$4,368.50 USD for LSTM and \$6,482.10 USD for TFDMGA. However, Fama-French 5-factor spanning regressions yield an estimated net strategy alpha of $\hat{\alpha} = -0.18\%$ per year ($p = 0.976, R^2 = 41.2\%$). Because net alpha is statistically zero, strategy returns

generated by these deep learning models represent dynamic compensation for systematic factor risk exposures rather than unpriced market inefficiency.

Keywords: Empirical Asset Pricing, Cross-Sectional Return Prediction, Deep Learning, Transformer Attention, Fama-French Spanning, Transaction Costs, Walk-Forward Validation.

يا رب زدني علما

“And say, ‘My Lord, increase me in knowledge.’ ”

, The Holy Qur’an, Surah Ta-Ha (20:114)

Acknowledgements

I express my sincere gratitude to my supervisor, **Prof. Giuseppina Chesini**. Her mentorship, academic guidance, and critical feedback throughout this research project shaped the empirical design and econometric focus of this Master's thesis.

I extend my thanks to the faculty and academic staff of the **Department of Economics at the University of Verona**. The coursework and quantitative training provided during my Master's studies laid the foundation for this work.

Moving from Mumbai to Verona to pursue my Master's degree has been a deeply rewarding journey. I am thankful to the University of Verona for offering an inspiring academic environment.

I acknowledge the financial support provided by **ESU Verona and the Borsa di Studio scholarship programme**, which enabled me to focus fully on my academic studies.

To my friends and classmates in Italy, thank you for your camaraderie, study sessions, and support during our time together in Verona.

I also recognize the open-source scientific community whose software libraries and computational frameworks made the empirical implementation of this study possible.

Above all, I thank my family. I am forever grateful to my father for his unwavering belief in my potential, his sacrifices, and his constant encouragement to pursue higher education. My gratitude also goes to my mother, my sister, and my extended family for their patience, warmth, and steady support.

Murtuza Rangwala

Master of Economics and Data Analysis
University of Verona

Contents

Abstract	i
Dedication	iii
Acknowledgements	iv
1 Introduction	1
1.1 Background and Economic Motivation	1
1.2 Research Questions and Core Hypotheses	2
1.3 Summary of Key Empirical Findings	3
1.4 Structure of the Thesis Manuscript	4
2 Literature Review	6
2.1 Introduction	6
2.2 Theoretical Asset Pricing Foundations	7
2.2.1 Equilibrium Foundations: The CAPM and Arbitrage Pricing Theory	7
2.2.2 The Fama-MacBeth Two-Pass Regression Framework	8
2.2.3 Multi-Factor Models: Fama-French, Carhart, and Pastor- Stambaugh	9
2.3 The Factor Zoo Controversy, Multiple Testing, and Shrinkage	11
2.3.1 The Proliferation of Anomalies and Cochrane's "Factor Zoo" . .	11
2.3.2 Multiple Testing Inflation and False Discovery Rates	12

2.3.3	High-Dimensional Regularized SDF Estimation	13
2.4	High-Dimensional SDF Estimation & Characteristic Interactions	14
2.4.1	Non-Parametric Characteristic Pruning via Group LASSO	14
2.4.2	Sparse Signals and Latent Factor Shrinkage	15
2.5	Machine Learning and Deep Learning Paradigms	16
2.5.1	Empirical Asset Pricing Benchmark: Gu, Kelly, and Xiu (2020)	16
2.5.2	Characteristics as Dynamic Risk Exposures: IPCA	17
2.5.3	Deep Learning and Non-Parametric SDF Estimation	17
2.5.4	Economic Frictions and Limits to Arbitrage	18
2.6	Financial Time-Series Econometrics & Inference	19
2.6.1	Heteroskedasticity and Autocorrelation Consistent (HAC) Infer- ence	19
2.6.2	Out-of-Sample Predictive Accuracy: Diebold-Mariano and HLN Corrections	20
2.6.3	Data Mining Diagnostics: Reality Check and Superior Predictive Ability	21
2.7	Financial ML Protocols & Overfitting Mitigation	21
2.7.1	Methodological Pitfalls in Financial Machine Learning	21
2.7.2	Validation Protocols	22
2.7.3	Microstructure Frictions & Transaction Cost Taxonomy	22
2.8	Synthesis of Core Literature Benchmarks	23
2.9	Summary	23
3	Data & Methodology	24
3.1	Data Pipeline and Sample Construction	24

3.1.1	Point-in-Time Bloomberg Filing Date Alignment	25
3.1.2	Survivorship-Bias-Free Dynamic Panel	25
3.2	Feature Space Taxonomy and Selection	26
3.2.1	Economic Rationale for Feature Set Size (59 Variables vs. 120+ Zoo Features)	26
3.3	Feature Normalization and Selection Protocol	27
3.4	Baseline Econometric & Machine Learning Formulations	29
3.4.1	Fama-MacBeth 2-Pass Linear Regression & Newey-West HAC Estimation	30
3.4.2	Regularized Linear Models: ElasticNet, LASSO, & Ridge	34
3.4.3	Non-Linear Tree Ensembles: Random Forest & XGBoost	37
3.4.4	Baseline Sequence Encoder: PyTorch LSTM	41
3.5	Custom Architecture: TFDMGA Network	42
3.5.1	Temporal Feature Encoding via Causal 1D Dilated TCN	43
3.5.2	Sequential Directed Ring Attention Cascade	44
3.5.3	3-Way Macro Dynamic Gating Network	45
3.5.4	Residual Fusion Block & Pre-LN Transformer Encoder	47
3.5.5	Multi-Task Multi-Head Objective Function	48
3.6	Cross-Validation & Walk-Forward Evaluation Scheme	49
3.7	Hyperparameter Search Space & Training Setup	50
3.8	Computational Complexity and Resource Profile	51
4	Empirical Results & Backtesting	52
4.1	Introduction	52
4.2	Contemporaneous Fama-MacBeth OLS Baseline	52

4.3	Mathematical Derivation & Empirical Reconciliation of LASSO Shrinkage Collapse	53
4.4	SHAP Feature Interpretability Analysis	54
4.5	Out-of-Sample Predictive Evaluation	56
4.6	Ablation Study of Neural Architecture Components	57
4.7	Robustness Checks and Sensitivity Analysis	57
4.8	Portfolio Construction & Execution Friction	58
4.9	Transaction Cost Sensitivity & \$1,000 USD Account Compounding . . .	58
4.10	Fama-French 5-Factor Spanning Regressions	60
5	Conclusion & Future Extensions	62
5.1	Summary of Research Findings	62
5.2	Theoretical and Practical Implications	63
5.3	Limitations of the Study	64
5.4	Future Research Extensions	64
	Bibliography	66

List of Figures

3.1	Bloomberg Point-in-Time Data Pipeline and Staleness Truncation Flowchart	26
3.2	Variable Taxonomy and Feature Group Modality Split (53 ML Inputs + 6 Factor Betas = 59 Total)	27
3.3	Spearman Rank Correlation Heatmap of Selected 53 ML Features	28
3.4	Stage 3 In-Sample Permutation Feature Importance Rankings across Modalities	29
3.5	Architecture Diagram of the Temporal Fusion Deep Multimodal Gated Attention (TFDMGA) Network	42
3.6	5-Fold Expanding-Window Walk-Forward Validation Protocol (2015–2024)	49
4.1	Global SHAP Mean Absolute Feature Importance Rankings (21-Day Horizon)	54
4.2	SHAP Summary Beeswarm Plot Illustrating Feature Value Impact on Model Predictions	55
4.3	Quintile Portfolio Sorting, 1-Day Execution Delay, and Turnover Rebalancing Flowchart	58
4.4	Out-of-Sample Cumulative Return Trajectories (Net 10 bps, 2020–2024 Folds)	59
4.5	Rolling 126-Day Annualized Sharpe Ratios across Model Architectures	59
4.6	Impact of 2:1 Take-Profit/Stop-Loss Risk Overlay (+4% / -2%) on Portfolio Returns	59

4.7	Maximum Drawdown Mitigation During Macro Stress Periods (COVID 2020 & Fed Hikes 2022)	60
-----	--	----

List of Tables

2.1	Comprehensive Synthesis of Core Literature Benchmarks	23
3.1	Hyperparameter Search Space and Optimal Tuned Parameters Across Models	50
3.2	Computational Complexity, Parameter Count, and Resource Execution Profile	51
4.1	Fama-MacBeth Daily Cross-Sectional Regression Risk Premia (2015– 2024)	52
4.2	Out-of-Sample Predictive Performance Metrics (2020–2024 Test Folds – Master Panel 53 Features)	56
4.3	Systematic Component Ablation Study of the TFDMGA Network . . .	57
4.4	Predictive Accuracy and Account Compounding Across Prediction Hori- zons and Fees	58
4.5	\$1,000 USD Account Growth Under Transaction Cost Scenarios (2020– 2024)	58
4.6	Fama-French 5-Factor Spanning Regressions with Newey-West HAC Errors	60

1. Introduction

1.1 Background and Economic Motivation

Why do financial asset returns vary across stocks, and how effectively do market prices incorporate available information? Answering this fundamental question lies at the heart of empirical asset pricing (Fama, 1970; Fama and French, 1993). For decades, academic researchers modeled expected equity returns through linear factor frameworks, evolving from Markowitz’s (1952) mean-variance framework to the Capital Asset Pricing Model (CAPM) of Sharpe (1964) and Lintner (1965), Ross’s (1976) Arbitrage Pricing Theory, Carhart’s (1997) four-factor momentum specification, Pástor and Stambaugh’s (2003) liquidity model, Hou, Xue, and Zhang’s (2015) q -factor framework, and Fama and French’s (2015) five-factor model. These classic models share a core assumption: stock returns are governed linearly by a small set of static risk exposures.

Real-world financial markets, however, display non-linearities, heavy-tailed return distributions, and dynamic regime shifts. Firm-level characteristics interact in nonlinear ways, for instance, valuation metrics behave differently in market drawdowns compared to economic expansions (Gu et al., 2020; Kozak et al., 2020). Standard Ordinary Least Squares (OLS) regressions face structural hurdles in high dimensional panels, risking omitted variable bias when specifications are too simple, or variance inflation when dozens of firm characteristics are included simultaneously.

Machine learning (ML) and deep sequence architectures provide a flexible framework for modeling nonlinear characteristic interactions and temporal patterns in equity markets (Gu et al., 2020, 2021; Chen et al., 2024; Kozak et al., 2020; Freyberger et al., 2020; Kelly et al., 2019). Techniques ranging from regularized linear models (Tibshirani, 1996; Zou and Hastie, 2005) and decision tree ensembles (Breiman, 2001; Friedman, 2001; Chen and Guestrin, 2016) to temporal convolutional networks (Bai et al., 2018) and multimodal Transformer architectures (Vaswani et al., 2017; Lim et al., 2021) allow researchers to extract predictive signals directly from raw panel data.

Yet applying machine learning to financial asset pricing differs from standard computer vision or language modeling. Financial return series are characterized by low signal-to-noise ratios, regime shifts, and structural breaks (López de Prado, 2018, 2020). Naive applications risk lookahead bias, survivorship bias, multiple testing inflation (Harvey et al., 2016; Harvey and Liu, 2020; Hou et al., 2020), and unmodeled transaction cost drag (Novy-Marx and Velikov, 2016; McLean and Pontiff, 2016; Avramov et al., 2023). Without out-of-sample walk-forward validation and trading friction modeling, apparent predictive gains can quickly disappear in live trading.

In this thesis, I evaluate machine learning models for cross-sectional return prediction across S&P 500 equities over the 2015–2024 period, addressing these methodological challenges directly.

1.2 Research Questions and Core Hypotheses

To evaluate the statistical predictability and economic viability of machine learning models, I frame my investigation around four central research questions:

Research Question 1 (Statistical Predictability): Do nonlinear machine learning models (Random Forest, XGBoost, LSTM, TFDMDGA) achieve statistically significant out-of-sample return predictability compared to linear econometric baselines (Fama-MacBeth OLS, LASSO, ElasticNet)?

Hypothesis 1: Nonlinear models, especially sequence-aware deep neural networks (Bai et al., 2018; Chen et al., 2024), outperform linear baselines in out-of-sample Information Coefficient (IC), Information Ratio ($ICIR$), and Directional Accuracy by capturing nonparametric characteristic interactions and temporal dependencies.

Research Question 2 (Multi-Modal Information Integration): Does integrating technical indicators, quarterly Bloomberg accounting ratios, macroeconomic state variables, and news/social sentiment within a dynamic gating architecture improve prediction accuracy over single-modality models?

Hypothesis 2: Multimodal feature integration (Cong et al., 2021; Lim et al., 2021) allows dynamic attention mechanisms to reweight feature modalities as macroeconomic regimes shift, yielding higher out-of-sample correlation and signal stability.

Research Question 3 (Economic Significance & Transaction Friction): Can out-of-sample return forecasts generate profitable trading strategies after accounting for transaction frictions (10 bps round-trip fees, execution delays, turnover, and stop-loss risk management)?

Hypothesis 3: While daily portfolio rebalancing incurs cost drag under 10 bps transaction fees (Novy-Marx and Velikov, 2016; Avramov et al., 2023), deep sequence models paired with dynamic take-profit/stop-loss overlays help control drawdowns and compound capital over time.

Research Question 4 (Econometric Factor Spanning & Market Efficiency): Are out-of-sample returns produced by deep learning models indicative of unpriced market alpha ($\alpha > 0$), or are they fully spanned by systematic risk factors ($\alpha = 0$) within the Fama-French 5-Factor framework?

Hypothesis 4: Out-of-sample strategy returns are explained by systematic factor risk exposures ($Mkt - RF, SMB, HML, RMW, CMA$) (Fama and French, 2015; Kelly et al., 2019), indicating that machine learning models act as dynamic factor allocation engines rather than discoverers of unpriced arbitrage.

1.3 Summary of Key Empirical Findings

Using a 5-fold expanding-window walk-forward validation scheme over the 2020–2024 out-of-sample period (detailed in Chapter 3 and Chapter 4), my main findings are:

1. **Performance in Return Rank Forecasting:** The custom Temporal Fusion Deep Multimodal Gated Attention (**TFDMGA**) network achieves the highest predictive

accuracy among all evaluated models, recording an out-of-sample daily Information Coefficient of $IC = +0.0348$, an Information Ratio of $ICIR = 3.12$, a ROC AUC of 0.6120, and a directional accuracy of 56.84%.

2. **Statistical Comparison with Baseline Sequence Models:** A formal Diebold and Mariano (1995) forecast accuracy test with Harvey et al. (1997) small-sample corrections confirms that TFDMGA yields statistically significant performance gains over a baseline PyTorch LSTM model ($DM = 2.41, p = 0.016$), validating the structural benefits of crossmodal ring attention and dynamic macro gating.
3. **Capital Compounding Trajectory:** Under a realistic 10 bps round-trip transaction cost setting and a 1-day execution buffer (Novy-Marx and Velikov, 2016), an initial \$1,000 USD deposit in January 2020 grows to **\$3,120.99 USD** under the baseline LSTM Q_5 long strategy, and reaches **\$4,368.50 USD** with a 2:1 Take-Profit/Stop-Loss (TPSL) risk overlay. Under TFDMGA with the TPSL overlay, capital compounds to **\$6,482.10 USD** by December 2024 (see Table 4.5 in Chapter 4).
4. **Fama-French 5-Factor Spanning & Factor Timing:** Spanning regressions against the Fama and French (2015) 5-factor model yield an estimated net strategy alpha of $\hat{\alpha} = -0.18\%$ per year ($p = 0.976, R^2 = 41.2\%$). Because net alpha is statistically zero, strategy returns are explained by dynamic risk exposure to systematic factor risks, specifically operating profitability (RMW beta $= +0.512, t = +9.12$), consistent with market efficiency (Fama, 1970).

1.4 Structure of the Thesis Manuscript

The remainder of this thesis is structured across four consolidated chapters:

- **Chapter 2 (Literature Review)** surveys academic literature in empirical asset pricing, factor models, high dimensional regularization, and financial machine learning.

- **Chapter 3 (Data & Methodology)** details the Bloomberg point-in-time data pipeline, constituent panel construction, 59-variable feature taxonomy, baseline model formulations, and the custom TFDMGA network architecture.
- **Chapter 4 (Empirical Results & Backtesting)** presents daily Fama-MacBeth baselines, LASSO shrinkage derivations, SHAP interpretability analysis, out-of-sample predictive metrics, ablation results, trading backtests under transaction fees, and Fama-French spanning regressions.
- **Chapter 5 (Conclusion & Future Research Extensions)** summarizes the main findings, discusses theoretical and practical implications, addresses study limitations, and outlines future research extensions.

2. Literature Review

2.1 Introduction

Understanding stock return predictability requires tracing the evolution of empirical asset pricing over the past six decades. Early academic frameworks established that asset prices in competitive markets reflect available information (Fama, 1970), modeling expected excess returns as linear functions of systematic risk exposures (Sharpe, 1964; Lintner, 1965). Subsequent econometric advances operationalized these principles through multifactor cross-sectional regression models (Fama and MacBeth, 1973; Ross, 1976; Fama and French, 1993, 2015).

However, the rapid expansion of financial data led to what Cochrane (2011) famously termed the “factor zoo”, a proliferation of over 300 candidate anomaly variables claiming to generate alpha. In response, modern financial econometrics has developed along two primary paths: (i) statistical auditing of data-mining inflation and false discovery rates (Harvey et al., 2016; Harvey and Liu, 2020; Hou et al., 2020), and (ii) high dimensional regularized estimation and nonparametric machine learning (Kozak et al., 2020; Freyberger et al., 2020; Gu et al., 2020; Chen et al., 2024).

In this chapter, I synthesize the asset pricing literature around six key thematic pillars that ground the theoretical framework, model architectures, feature selection, and statistical inference used in this thesis:

1. **Theoretical Asset Pricing Foundations:** From the single factor equilibrium of CAPM and APT to the multifactor models of Fama-French, Carhart, and Pastor-Stambaugh.
2. **The Factor Zoo, Multiple Testing, and Shrinkage:** Multiple testing inflation, false discovery rates, and high dimensional regularized SDF estimation.

3. **High-Dimensional SDF Estimation & Characteristic Interactions:** Nonparametric characteristic selection, latent factor models, and dynamic risk loadings.
4. **Machine Learning & Deep Learning Paradigms:** Nonlinear empirical asset pricing, instrumented principal components, and deep sequence SDF models.
5. **Financial Time-Series Econometrics & Inference:** HAC covariance estimation, out-of-sample forecast accuracy tests, and data-mining diagnostics.
6. **Financial ML Protocols & Overfitting Mitigation:** Backtesting rules, expanding walk-forward splits, microstructural frictions, and transaction cost frameworks.

2.2 Theoretical Asset Pricing Foundations

2.2.1 Equilibrium Foundations: The CAPM and Arbitrage Pricing Theory

Modern quantitative finance originates with the mean-variance portfolio theory of Markowitz (1952) and the Capital Asset Pricing Model (CAPM) developed by Sharpe (1964) and Lintner (1965). Under the assumptions of homogeneous investor expectations, frictionless markets, and quadratic utility or normal return distributions, investors hold identical portfolios of risky assets weighted by market capitalizations. As a result, the market portfolio R_m is mean-variance efficient, and an asset's expected return is linearly related to its covariance with the aggregate market portfolio:

$$\mathbb{E}[R_{i,t}] - R_{f,t} = \beta_{i,m} (\mathbb{E}[R_{m,t}] - R_{f,t}) \quad (2.1)$$

where the market beta $\beta_{i,m} = \frac{\text{Cov}(R_{i,t}, R_{m,t})}{\text{Var}(R_{m,t})}$ measures systemic market risk. Black (1972) relaxed the assumption of risk-free borrowing, introducing the zero-beta CAPM and showing that the security market line remains linear but exhibits a higher intercept and flatter slope when borrowing constraints bind.

Recognizing the strict equilibrium assumptions of the CAPM, Ross (1976) formulated the Arbitrage Pricing Theory (APT). Based on the absence of asymptotic

arbitrage rather than general equilibrium, APT posits that asset returns follow a K -factor linear generative structure:

$$R_{i,t} = \mathbb{E}[R_{i,t}] + \sum_{k=1}^K \beta_{i,k} f_{k,t} + \epsilon_{i,t} \quad (2.2)$$

where $f_{k,t}$ represent mean-zero common macroeconomic risk factors, $\beta_{i,k}$ are factor loadings, and $\epsilon_{i,t}$ is an idiosyncratic error satisfying $\text{Cov}(\epsilon_{i,t}, \epsilon_{j,t}) = 0$ for all $i \neq j$. By applying linear pricing space arguments in infinite-dimensional Hilbert spaces, Ross showed that expected returns are linearly spanned by factor risk premia λ_k :

$$\mathbb{E}[R_{i,t}] - R_{f,t} = \sum_{k=1}^K \beta_{i,k} \lambda_k \quad (2.3)$$

2.2.2 The Fama-MacBeth Two-Pass Regression Framework

To empirically test whether factor loadings explain the cross-section of expected returns, Fama and MacBeth (1973) established the classic two-pass regression methodology. In the first pass, time series regressions are conducted for each asset i over an estimation window T_1 to estimate factor loadings $\hat{\beta}_{i,k}$:

$$R_{i,t} - R_{f,t} = \alpha_i + \sum_{k=1}^K \beta_{i,k} f_{k,t} + \epsilon_{i,t}, \quad t = 1, \dots, T_1 \quad (2.4)$$

In the second pass, cross-sectional regressions are executed at each time step t over an evaluation window T_2 using the estimated factor loadings as explanatory variables:

$$R_{i,t} - R_{f,t} = \gamma_{0,t} + \sum_{k=1}^K \hat{\beta}_{i,k} \gamma_{k,t} + u_{i,t}, \quad i = 1, \dots, N \quad (2.5)$$

The point estimates of factor risk premia $\bar{\gamma}_k$ and their standard errors are obtained by averaging the cross-sectional slope coefficients across time:

$$\bar{\gamma}_k = \frac{1}{T_2} \sum_{t=1}^{T_2} \hat{\gamma}_{k,t}, \quad \text{SE}(\bar{\gamma}_k) = \sqrt{\frac{1}{T_2(T_2 - 1)} \sum_{t=1}^{T_2} (\hat{\gamma}_{k,t} - \bar{\gamma}_k)^2} \quad (2.6)$$

While standard Fama-MacBeth standard errors assume cross-sectional independence and zero serial correlation, modern implementations incorporate Heteroskedasticity and Autocorrelation Consistent (HAC) corrections (Newey and West, 1987) and Shanken (1992) corrections for generated regressor error in $\hat{\beta}_{i,k}$.

2.2.3 Multi-Factor Models: Fama-French, Carhart, and Pastor-Stambaugh

Empirical challenges to the single factor CAPM increased during the late 1980s and early 1990s as cross-sectional anomalies emerged. Fama and French (1992) showed that firm size (market capitalization) and book-to-market equity ratio (B/M) absorbed much of the cross-sectional explanatory power of market beta. To formalize these empirical findings, Fama and French (1993) introduced the Fama-French Three-Factor Model:

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2}SMB_t + \beta_{i,3}HML_t + e_{i,t} \quad (2.7)$$

where SMB_t (Small Minus Big) captures the size effect, and HML_t (High Minus Low) captures the value effect.

Observing that the three-factor model failed to price the cross-sectional price momentum phenomenon identified by Jegadeesh and Titman (1993), Carhart (1997) expanded the framework into a Four-Factor Model by adding a momentum factor (WML_t , Winners Minus Losers):

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2}SMB_t + \beta_{i,3}HML_t + \beta_{i,4}WML_t + e_{i,t} \quad (2.8)$$

In parallel, Pástor and Stambaugh (2003) showed that market-wide liquidity is an un-diversifiable risk factor. Stocks with high sensitivity to aggregate market liquidity fluctuations command higher expected returns to compensate investors for illiquidity shocks during market downturns.

To address profitability and investment anomalies, Fama and French (2015) extended their three-factor model to a Five-Factor Asset Pricing Model:

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2} SMB_t + \beta_{i,3} HML_t + \beta_{i,4} RMW_t + \beta_{i,5} CMA_t + e_{i,t} \quad (2.9)$$

where RMW_t (Robust Minus Weak) measures the profit margin return differential, and CMA_t (Conservative Minus Aggressive) captures investment intensity differentials. Grounded theoretically in the dividend discount model, Fama and French (2015) showed that holding book-to-market fixed, higher expected earnings imply higher expected returns, while higher expected investment rates lower expected returns. Concurrently, Hou et al. (2015) established the q -factor model ($Mkt, ME, I/A, ROE$), showing that investment and profitability factors derived from q -theory span a wide range of cross-sectional anomalies.

2.3 The Factor Zoo Controversy, Multiple Testing, and Shrinkage

2.3.1 The Proliferation of Anomalies and Cochrane's "Factor Zoo"

By the early 2010s, academic literature had documented over 300 firm characteristics claiming to predict stock returns. In his American Finance Association presidential address, Cochrane (2011) famously diagnosed this proliferation as the “factor zoo.” Cochrane noted that while early theoretical models assumed a small number of macro risk factors span expected returns, the literature had created a large collection of firm characteristics without establishing whether candidate factors were incremental to existing benchmarks or merely duplicate representations of underlying noise.

This issue was compounded by post-publication return decay. McLean and Pontiff (2016) examined 97 published cross-sectional anomalies and documented that out-of-sample anomaly portfolio returns decline by an average of 26% after academic publication and up to 58% after original sample periods. This decay stems from a combination of market learning (arbitrage capital exploiting newly disclosed anomalies) and data-mining biases in published literature.

2.3.2 Multiple Testing Inflation and False Discovery Rates

The main statistical cause of factor zoo inflation was analyzed by Harvey et al. (2016). Traditional empirical asset pricing evaluated candidate factors using single-hypothesis Student's t -tests with standard significance thresholds ($|t| > 1.96$, corresponding to $\alpha = 0.05$). However, when hundreds of candidate signals are evaluated on the same datasets, standard significance cutoffs produce high false discovery rates due to multiple hypothesis testing.

Harvey et al. (2016) evaluated over 300 published anomaly factors under multiple testing adjustments (including Bonferroni corrections and Holm-Bonferroni procedures). They showed that to maintain a true False Discovery Rate (FDR) below 5%, the minimum t -statistic cutoff for newly proposed asset pricing factors should be elevated to at least:

$$|t| > 3.0 \quad (p\text{-value} < 0.0027) \quad (2.10)$$

Extending this framework, Harvey and Liu (2020) introduced Bayesian FDR frameworks that account for prior belief distributions on true alpha. In a large replication audit, Hou et al. (2020) re-examined 452 cross-sectional anomalies under microcap exclusions and value-weighted portfolio rules. They showed that 65% of published anomalies fail to achieve statistical significance at standard $|t| > 1.96$ thresholds, and over 85% fail under the $|t| > 3.0$ threshold proposed by Harvey et al. (2016).

2.3.3 High-Dimensional Regularized SDF Estimation

To address the factor zoo problem structurally, econometricians re-framed asset pricing as a high dimensional regularized Stochastic Discount Factor (SDF) estimation problem. In a stochastic discount factor framework, an asset pricing model specifies a scalar pricing kernel M_{t+1} such that for all excess returns $R_{i,t+1}$:

$$\mathbb{E}_t [M_{t+1} R_{i,t+1}] = 0 \quad (2.11)$$

When the SDF is parameterized as a linear combination of a high dimensional vector of K candidate characteristics or factor returns $\mathbf{z}_{t+1}$, i.e., $M_{t+1} = 1 - \mathbf{b}^T(\mathbf{z}_{t+1} - \mathbb{E}[\mathbf{z}_{t+1}])$, unregularized sample estimators fail due to multicollinearity and $K \approx T$ sample noise.

Kozak et al. (2020) showed that shrinking the factor weighting vector $\mathbf{b}$ via combined L_1 (LASSO) and L_2 (Ridge) regularization yields robust, out-of-sample optimal SDF estimators. The regularized SDF objective function is formulated as:

$$\min_{\mathbf{b}} \left\{ \frac{1}{T} \sum_{t=1}^T (1 - \mathbf{b}^T \mathbf{z}_t)^2 + \gamma_1 \|\mathbf{b}\|_1 + \gamma_2 \|\mathbf{b}\|_2^2 \right\} \quad (2.12)$$

Kozak et al. (2020) showed that while unregularized models tend to overfit high-frequency sample noise, penalized SDF estimation extracts a low-dimensional principal component structure that spans expected returns while shrinking noise coefficients to zero.

2.4 High-Dimensional SDF Estimation & Characteristic Interactions

2.4.1 Non-Parametric Characteristic Pruning via Group LASSO

Parallel to regularized linear SDFs, Freyberger et al. (2020) developed a nonparametric model to systematically select which firm characteristics provide incremental predictive power for expected returns. Using nonparametric B-spline expansions and Group LASSO penalties, their estimator models expected returns as:

$$\mathbb{E}_t[R_{i,t+1}] = \sum_{j=1}^J g_j(c_{i,j,t}) + \epsilon_{i,t+1} \quad (2.13)$$

where $c_{i,j,t}$ represents the rank-transformed characteristic j for firm i at time t , and $g_j(\cdot)$ is a nonlinear continuous function. The Group LASSO penalty penalizes the vector of spline coefficients for each characteristic simultaneously:

$$\min_{\{g_j\}} \left\{ \frac{1}{NT} \sum_{i=1}^N \sum_{t=1}^T \left(R_{i,t+1} - \sum_{j=1}^J g_j(c_{i,j,t}) \right)^2 + \lambda \sum_{j=1}^J \sqrt{\dim(g_j)} \|g_j\|_2 \right\} \quad (2.14)$$

Freyberger et al. (2020) showed that out of more than 50 published characteristics, only a small subset (such as short term momentum, market capitalization, operating profitability, and asset growth) retain independent predictive power once non-linearities and characteristic redundancies are penalized.

2.4.2 Sparse Signals and Latent Factor Shrinkage

Chinco et al. (2019) examined sparse return predictability at high frequencies, showing that LASSO regression identifies transient, cross-asset predictive signals that standard time series estimators miss. Grounding model selection in econometric diagnostics, Lewellen et al. (2010) provided a critique of cross-sectional asset pricing tests. They showed that testing multifactor models on strong factor-structure portfolios (such as standard 25 Size-B/M sorted portfolios) produces artificially inflated cross-sectional R^2 values. They prescribed rigorous diagnostic standards: expanding test portfolios beyond Size-B/M, forcing theoretical factor risk premia to equal sample factor means, and reporting confidence intervals for cross-sectional R^2 .

Expanding structural factor selection, Feng et al. (2020) and Giglio and Xiu (2021) developed double-testing econometric frameworks combining LASSO factor selection with two-pass regressions. They showed that high dimensional regularization enables researchers to evaluate whether a newly proposed factor adds incremental explanatory power to the existing universe of 300+ factors without suffering from omitted variable bias or data mining contamination.

2.5 Machine Learning and Deep Learning Paradigms

2.5.1 Empirical Asset Pricing Benchmark: Gu, Kelly, and Xiu (2020)

The landmark study by Gu et al. (2020) established a primary theoretical and empirical benchmark for financial machine learning. Analyzing a panel of over 30,000 U.S. equities over 60 years (1957–2016) with 94 firm-level characteristics and 8 macroeconomic indicators, Gu, Kelly, and Xiu evaluated a range of machine learning architectures: Linear OLS, LASSO, ElasticNet, Principal Component Regression (PCR), Partial Least Squares (PLS), Random Forests, Gradient Boosted Decision Trees (GBDT), and multilayer Feedforward Neural Networks (MLP).

The nonlinear forecasting panel model is defined as:

$$R_{i,t+1} = g^*(\mathbf{x}_{i,t}) + \epsilon_{i,t+1}, \quad \mathbf{x}_{i,t} = \mathbf{c}_{i,t} \otimes \mathbf{d}_t \quad (2.15)$$

where $\mathbf{c}_{i,t}$ is a p -dimensional vector of firm characteristics, $\mathbf{d}_t$ is an m -dimensional vector of macro variables, and $\otimes$ denotes the Kronecker product creating interactive feature conditioning.

Gu et al. (2020) highlighted three main empirical findings:

1. **Non-Linear Model Performance:** Nonlinear models (GBDT, Neural Networks) outperform linear models in out-of-sample predictive R_{oos}^2 , doubling the annualized Sharpe ratios of long-short hedge portfolios.
2. **Importance of Interaction Terms:** A main source of machine learning gain comes from automatically capturing nonlinear characteristic interactions (such as the interaction of size with short term momentum and liquidity).
3. **Neural Network Architecture:** Shallow neural architectures with 1 to 3 hidden layers achieve better performance compared to deeper networks, as excessive depth can lead to overfitting on low signal-to-noise equity return data.

2.5.2 Characteristics as Dynamic Risk Exposures: IPCA

Theoretical motivation for why machine learning return forecasts succeed was provided by Kelly et al. (2019). Introducing Instrumented Principal Component Analysis (IPCA), Kelly, Pruitt, and Su modeled asset returns as:

$$R_{i,t+1} = \mathbf{z}_{i,t}^T \mathbf{\Gamma}_\beta \mathbf{f}_{t+1} + \epsilon_{i,t+1} \quad (2.16)$$

where $\mathbf{z}_{i,t}$ is a vector of observable firm characteristics, $\mathbf{\Gamma}_\beta$ is a mapping matrix connecting characteristics to dynamic factor loadings ($\beta_{i,t} = \mathbf{\Gamma}_\beta^T \mathbf{z}_{i,t}$), and $\mathbf{f}_{t+1}$ represents unobservable latent factor returns.

Kelly et al. (2019) showed that firm characteristics serve as proxies for time-varying factor risk exposures. Consequently, machine learning return predictors predicting $R_{i,t+1} = g(\mathbf{z}_{i,t})$ are implicitly estimating dynamic latent factor models with time-varying risk loadings.

2.5.3 Deep Learning and Non-Parametric SDF Estimation

Extending ML to advanced deep architectures, Chen et al. (2024) developed deep neural network models that estimate nonparametric asset pricing SDFs directly from high dimensional equity panels. By parameterizing factor loadings $\beta(\mathbf{c}_{i,t})$ as nonlinear functions of firm characteristics using Deep Neural Networks (DNNs) combined with Generative Adversarial Networks (GANs), Chen, Pelger, and Zhu showed that deep learning captures macroeconomic regime shifts and structural non-linearities that linear factor models may miss.

Similarly, Cong et al. (2021) incorporated multimodal deep learning architectures to process accounting signals alongside high dimensional textual information. To help explain the predictions of deep neural networks, Lundberg and Lee (2017) introduced

SHapley Additive exPlanations (SHAP), based on cooperative game theory. SHAP decomposes a model prediction $f(\mathbf{x})$ into additive feature contributions:

$$f(\mathbf{x}) = \phi_0 + \sum_{j=1}^M \phi_j \mathbf{x}_j \quad (2.17)$$

providing transparent economic interpretation of nonlinear asset pricing models.

2.5.4 Economic Frictions and Limits to Arbitrage

While machine learning models achieve notable statistical R_{oos}^2 , Avramov et al. (2023) investigated whether these return forecast gains survive real-world economic frictions. Avramov, Cheng, and Metzker showed that machine learning return predictability is heavily concentrated in illiquid, microcap stocks with high distress risk, high transaction costs, and borrowing constraints. When portfolios exclude microcaps or account for realistic execution costs, out-of-sample profits decline significantly. This highlights the importance of integrating microstructural frictions into financial machine learning evaluation.

2.6 Financial Time-Series Econometrics & Inference

2.6.1 Heteroskedasticity and Autocorrelation Consistent (HAC) Inference

Financial return time series exhibit conditional heteroskedasticity, volatility clustering, and serial dependence. Standard Ordinary Least Squares (OLS) covariance estimators understate parameter variances, yielding artificially inflated t -statistics. To provide valid statistical inference, Newey and West (1987) developed the Heteroskedasticity and Autocorrelation Consistent (HAC) covariance matrix estimator.

For a time series of regression residuals e_t and regressors $\mathbf{X}_t$, defining $\mathbf{u}_t = \mathbf{X}_t e_t$, the Newey-West estimator computes the sample covariance matrix $\hat{\Omega}_{\text{NW}}$ as:

$$\hat{\Omega}_{\text{NW}} = \hat{\Gamma}_0 + \sum_{j=1}^L w(j, L) \left(\hat{\Gamma}_j + \hat{\Gamma}_j^T \right) \quad (2.18)$$

where $\hat{\Gamma}_j = \frac{1}{T} \sum_{t=j+1}^T \mathbf{u}_t \mathbf{u}_{t-j}^T$ is the sample autocovariance matrix at lag j , and $w(j, L) = 1 - \frac{j}{L+1}$ represents the Bartlett kernel weighting function over a maximum lag length L . The Bartlett kernel ensures that $\hat{\Omega}_{\text{NW}}$ is positive semi-definite in finite samples.

2.6.2 Out-of-Sample Predictive Accuracy: Diebold-Mariano and HLN Corrections

When comparing out-of-sample forecast accuracy between two competing forecasting models (such as a deep sequence network versus a baseline linear factor model), standard independent sample t -tests are invalid due to forecast error autocorrelation. Diebold and Mariano (1995) introduced a distribution-free statistic for testing the null hypothesis of equal predictive accuracy.

Let $e_{1,t}$ and $e_{2,t}$ denote forecast errors from Model 1 and Model 2 at time t . The loss differential sequence is defined under a loss function $L(\cdot)$ (such as Mean Squared Error):

$$d_t = L(e_{1,t}) - L(e_{2,t}) \quad (2.19)$$

The null hypothesis of equal forecast accuracy is $H_0 : \mathbb{E}[d_t] = 0$. The Diebold-Mariano (DM) test statistic is given by:

$$DM = \frac{\bar{d}}{\sqrt{\frac{\widehat{V}(\bar{d})}{T}}} \xrightarrow{d} \mathcal{N}(0, 1) \quad (2.20)$$

where $\bar{d} = \frac{1}{T} \sum_{t=1}^T d_t$, and $\widehat{V}(\bar{d})$ is the HAC-adjusted asymptotic variance of the loss differential sequence using the Newey-West estimator.

To correct for small-sample size distortion in modest evaluation panels, Harvey et al. (1997) introduced the Harvey-Leybourne-Newbold (HLN) small-sample correction. For a forecast horizon h , the HLN-adjusted statistic is:

$$DM^* = DM \times \left[\frac{T + 1 - 2h + T^{-1}h(h-1)}{T} \right]^{1/2} \quad (2.21)$$

Critical values for DM^* are drawn from the Student's t -distribution with $(T-1)$ degrees of freedom, providing robust inference in empirical financial panels.

2.6.3 Data Mining Diagnostics: Reality Check and Superior Predictive Ability

When searching across multiple model hyperparameter configurations or architecture variants, researchers risk selecting a model that performs well out-of-sample purely by chance. To control for data mining, White (2000) introduced the Reality Check for Data Mining. White formulated a bootstrap procedure over the joint distribution of benchmark and candidate model forecasts to test whether the best performing candidate model genuinely outperforms the benchmark.

Recognizing that White’s Reality Check can be conservative when poor candidate models are included, Hansen (2005) refined the framework into the Superior Predictive Ability (SPA) test. Hansen introduced a studentized test statistic with sample-dependent variance adjustments:

$$T_k^{\text{SPA}} = \max_{k=1,\dots,K} \frac{\sqrt{T} \bar{d}_k}{\hat{\omega}_k} \quad (2.22)$$

where $\hat{\omega}_k^2$ is a consistent estimator of the asymptotic variance of $\sqrt{T} \bar{d}_k$. Hansen showed that the SPA test exhibits higher power and robustness against uninformative candidate models, serving as a useful diagnostic for quantitative model comparison.

2.7 Financial ML Protocols & Overfitting Mitigation

2.7.1 Methodological Pitfalls in Financial Machine Learning

In financial machine learning, standard statistical cross-validation techniques (such as standard randomized K -fold CV) can fail. As detailed by López de Prado (2018), financial time series data violate the independent and identically distributed (i.i.d.) assumption. Applying naive cross-validation induces several methodological problems:

1. **Lookahead Leakage:** Using future information (such as future target returns or lookahead normalized features) during training.
2. **Serial Memory Contamination:** Overlapping target return windows (such as multi-day cumulative returns) allow training samples to overlap with test samples, leaking labels across folds.

3. **Backtest Overfitting:** Iteratively tuning model parameters directly on test set Sharpe ratios until an illusion of profitability is produced.

2.7.2 Validation Protocols

To ensure backtesting integrity, López de Prado (2018) established validation protocols:

- **Expanding Walk-Forward Validation:** Models are trained on historical data $t \in [1, T_{\text{train}}]$ and evaluated on subsequent out-of-sample periods $t \in [T_{\text{train}} + 1, T_{\text{test}}]$. The training window expands iteratively over time.
- **Combinatorial Purged Cross-Validation (CPCV):** When non-sequential cross-validation is conducted, training folds adjacent to test folds are *purged* by removing samples whose label windows overlap with test windows. In addition, an *embargo* period is appended after test sets to eliminate serial correlation leakage.
- **Point-in-Time Data Alignment:** Financial fundamentals and accounting metrics are indexed by their actual public SEC disclosure dates (such as 10-K/10-Q filing dates) rather than fiscal period end dates.

2.7.3 Microstructure Frictions & Transaction Cost Taxonomy

High-frequency predictive signals can decay rapidly once realistic transaction costs and market microstructure frictions are accounted for (Novy-Marx and Velikov, 2016; Frazzini and Pedersen, 2014). Novy-Marx and Velikov (2016) show that many published anomaly strategies disappear after deducting effective bid-ask spreads, market impact costs, and turnover frictions.

Second, execution realism mandates an execution delay: trading signals generated at market close on day t cannot be executed until day $t + 1$ close or open. In addition, leverage constraints and margin costs (Frazzini and Pedersen, 2014) restrict short-selling capacity on distressed equities, reinforcing the necessity of evaluating long-only and cost-adjusted long-short performance.

2.8 Synthesis of Core Literature Benchmarks

Table 2.1 provides a synthesis of the literature that anchors the theoretical framing, model architecture, feature engineering, econometric estimation, and validation of this thesis.

Table 2.1: Comprehensive Synthesis of Core Literature Benchmarks

Thematic Domain	Key Literature Citation	Core Theoretical / Econometric Finding	Implementation in Thesis
Asset Pricing Foundations	Sharpe (1964), Lintner (1965)	Single factor CAPM equilibrium; market beta linear pricing	Baseline risk benchmark
	Fama and MacBeth (1973)	Two-pass cross-sectional regression methodology	Baseline linear factor risk premia
	Ross (1976)	Arbitrage Pricing Theory (APT); multifactor space	Linear multifactor baseline
	Fama and French (1993), Fama and French (2015)	3-Factor and 5-Factor models (<i>Mkt, SMB, HML, RMW, CMA</i>)	Econometric spanning regressions
	Carhart (1997), Pastor and Stambaugh (2003)	Price momentum (<i>WML</i>) and aggregate market liquidity factors	Extended multifactor benchmark
Factor Zoo & Shrinkage	Cochrane (2011)	Factor zoo diagnosis; lack of unified factor structure	High dimensional model motivation
	Harvey et al. (2016), Harvey and Liu (2020)	Multiple testing bias; minimum hurdle rate $ t > 3.0$	Feature auditing threshold
	Hou et al. (2020), McLean and Pontiff (2016)	65%+ anomaly replication failure; post-publication decay	Microcap exclusions & value weighting
	Kozak et al. (2020)	Regularized L_1/L_2 SDF shrinkage in high dimensions	Penalized Ridge/ElasticNet baselines
Machine Learning & Deep Learning	Freyberger et al. (2020)	Nonparametric Group LASSO characteristic pruning	Rank-normalized feature selection
	Gu et al. (2020)	Nonlinear ML performance; interaction terms; 1–3 layer networks	Experimental lineup design
	Kelly et al. (2019)	Instrumented PCA (IPCA); characteristics as dynamic betas	Latent dynamic factor framing
	Chen et al. (2024)	Deep learning nonparametric SDF estimation via GANs	Sequence model architecture
Time-Series Econometrics	Avramov et al. (2023)	ML gains concentrated in illiquid hard-to-arbitrage stocks	Size-stratified empirical evaluation
	Lundberg and Lee (2017)	SHAP additive feature attribution for nonlinear models	Model interpretability framework
	Newey and West (1987)	HAC robust standard errors; Bartlett kernel weighting	Fama-MacBeth HAC t-statistics
Financial ML Validation	Diebold and Mariano (1995), Harvey et al. (1997)	Out-of-sample forecast accuracy test & HLN correction	Model forecast comparison tests
	White (2000), Hansen (2005)	Reality Check and Superior Predictive Ability (SPA) tests	Hyperparameter data-mining audit
	López de Prado (2018)	Financial ML protocols; purging, embargoing, walk-forward	5-Fold expanding walk-forward split
	Novy-Marx and Velikov (2016)	Transaction cost taxonomy; effective spreads and impact	10 bps friction drag evaluation
	Frazzini and Pedersen (2014)	Betting against beta; leverage constraints and margin drag	Execution delay & shorting rules

2.9 Summary

In conclusion, the evolution of empirical asset pricing shows a clear transition from simple linear factor models toward regularized high dimensional estimation and nonparametric deep sequence architectures. By synthesizing the theoretical insights of factor pricing, the statistical rigor of multiple testing adjustments, the predictive capability of deep sequence networks, and empirical validation protocols, this thesis builds a framework for cross-sectional stock return prediction.

3. Data & Methodology

3.1 Data Pipeline and Sample Construction

My empirical sample comprises daily stock prices, trading volumes, and quarterly financial accounting metrics for active and historical S&P 500 constituent equities from December 31, 2014, through December 31, 2024. The active evaluation window spans January 1, 2015, to December 31, 2024, yielding 2,516 trading days and over 1.25 million stock-day observations (Gu et al., 2020).

Daily closing prices, dividend payments, stock splits, and trading volumes are ingested from Yahoo Finance and cross-verified against Bloomberg Terminal records. Daily risk-free rates ($R_{f,t}$) along with daily Fama-French 5-Factor returns ($Mkt - RF_t, SMB_t, HML_t, RMW_t, CMA_t$) and Carhart Momentum (MOM_t) are retrieved directly from the Ken French Data Library (Fama and French, 1993, 2015; Carhart, 1997).

3.1.1 Point-in-Time Bloomberg Filing Date Alignment

A common issue in empirical finance studies is lookahead bias caused by indexing corporate accounting metrics to fiscal period end-dates. In practice, financial statements for a fiscal quarter ending on December 31 (t_{fiscal}) are published weeks or months later during the 10-K or 10-Q SEC filing process (López de Prado, 2018).

To eliminate lookahead contamination, my data pipeline (`data_pipeline.py` in `src/`) ingests official Bloomberg publication timestamps (`earnings_announcement_date` and `filing_date`). On any given trading date t , a quarterly fundamental feature vector $\mathbf{x}_{i,t}^{\text{fund}}$ is updated to the newly reported accounting state Q_τ if and only if:

$$t \geq \text{Announcement Date}(Q_\tau) \quad (3.1)$$

Prior to $t = \text{Announcement Date}(Q_\tau)$, the model is strictly limited to using the previously announced fundamental vector $Q_{\tau-1}$.

In addition, I enforce a **Staleness Truncation Rule**: a fundamental observation remains active for at most 90 trading days (≈ 1 calendar quarter). If a firm fails to publish a new statement within 120 calendar days (≈ 90 trading days), the forward-fill is truncated and the corresponding fundamental fields are set to missing (NaN), preventing outdated accounting ratios from skewing cross-sectional ranks (Novy-Marx and Velikov, 2016).

3.1.2 Survivorship-Bias-Free Dynamic Panel

To guard against survivorship bias, my asset universe tracks historical point-in-time S&P 500 constituent rosters obtained from Bloomberg. Companies removed from the index, acquired, or delisted due to bankruptcy between 2015 and 2024 remain in the daily cross-section up to their final active trading date T_{exit} (McLean and Pontiff, 2016; Hou et al., 2020).

Figure 3.1 shows my point-in-time data alignment pipeline and staleness truncation flowchart.



Figure 3.1: Bloomberg Point-in-Time Data Pipeline and Staleness Truncation Flowchart

3.2 Feature Space Taxonomy and Selection

My empirical panel relies on **59 variables in total**, split between predictive model inputs and benchmark risk factors:

- **53 Predictive Machine Learning Inputs:** Fed directly into the machine learning and neural network architectures across four modalities: Technical (16), Fundamental (30), Sentiment (2), and Macro (5) (Gu et al., 2020; Cong et al., 2021).
- **6 Benchmark Risk Factor Betas:** Retained for baseline linear regressions and econometric risk adjustments, including the 5 Fama-French systematic factor betas ($Mkt - RF$, SMB , HML , RMW , CMA) and Momentum (MOM) (Fama and French, 2015; Carhart, 1997).

3.2.1 Economic Rationale for Feature Set Size (59 Variables vs. 120+ Zoo Features)

When designing the characteristic space, a key decision is whether to include hundreds of candidate anomaly signals or focus on a curated panel. I deliberately choose a 59-variable feature panel rather than expanding to 120 or 250 predictors.

Including vast numbers of firm characteristics in financial panels introduces practical and statistical problems. Many raw accounting metrics simply restate similar economic signals, such as minor variations of leverage ratios or momentum lookbacks, adding collinearity and noise rather than genuine predictive content (Freyberger et al., 2020; Kozak et al., 2020). In addition, as Harvey et al. (2016) observe, testing hundreds of features increases false-discovery rates. Restricting the predictor set to 53 ML inputs grounded in four economic domains (Technical, Fundamental, Sentiment, and Macro) ensures each variable represents an established economic mechanism (Pástor and Stambaugh, 2003; Novy-Marx, 2013; Cong et al., 2021; Fama and French, 2015). Finally, requiring hundreds of characteristics reduces cross-sectional sample coverage, as

missing fundamental items force the exclusion of valid stock-day observations (Avramov et al., 2023).

Figure 3.2 displays the feature group modality classification.



Figure 3.2: Variable Taxonomy and Feature Group Modality Split (53 ML Inputs + 6 Factor Betas = 59 Total)

3.3 Feature Normalization and Selection Protocol

To maintain uniform scaling and stationarity, raw characteristic variables are converted each day into **Cross-Sectional Ranks** mapped to the unit interval $[0, 1]$ (Freyberger et al., 2020; Gu et al., 2020):

$$r_{i,t,j} = \frac{\text{Rank}(X_{i,t,j}) - 1}{N_t - 1} \in [0, 1] \quad (3.2)$$

Rank-scaling protects models from extreme financial outliers, standardizes variables across different accounting units, and keeps distributions comparable across changing market environments.

Importantly, cross-sectional rank normalization in Equation (3.2) is applied strictly to firm-level characteristics (Technical and Fundamental features). Macroeconomic state variables (x_{macro}) and news/social sentiment indicators are **not** rank-normalized cross-sectionally across stocks on day t — doing so would reduce their cross-sectional variance to zero, as macro variables are uniform across all assets on a given trading day. Instead, macro state variables are normalized temporally across time or processed directly by the 3-Way Macro Dynamic Gating network.

To guarantee zero lookahead leakage and zero overnight gap contamination, return targets $y_{i,t+1}$ in training and evaluation are strictly computed from $P_{i,t+1,\text{open}}$ to $P_{i,t+2,\text{open}}$ (or $P_{i,t+1,\text{close}}$ to $P_{i,t+2,\text{close}}$ following execution at $t + 1$), matching the 1-day execution buffer delay ($\text{shift}(-2)$).

To avoid data leakage, I perform feature selection (`select_features.py` in `scripts/`) using a **three-stage protocol applied exclusively to pre-2020 training data** (146,959 observations) (López de Prado, 2018). The first stage removes features with missing rates above 10%. The second stage prunes redundant variables exhibiting pairwise Spearman correlations $|\rho_{j,k}| > 0.85$. In the third stage, an in-sample Random Forest regressor (Breiman, 2001) ranks remaining candidate variables, selecting the top 53 predictive features.

Figure 3.3 presents the correlation matrix of selected features, and Figure 3.4 shows the Random Forest importance rankings.




Figure 3.3: Spearman Rank Correlation Heatmap of Selected 53 ML Features

Figure 3.4 reports the in-sample permutation feature importances calculated strictly on pre-2020 training data to prevent out-of-sample lookahead leakage. A classic methodological issue in tree-based feature selection is that raw Gini impurity measures over-attribute importance to continuous, un-normalized variables (Strobl et al., 2007). Evaluating features via out-of-sample permutation importance and cross-sectional rank normalization eliminates this continuous cardinality bias. Technical momentum indicators (vol_{126d} , mom_{21d} , rsi_{14d}) and fundamental metrics (pe_{ratio} , roi , quality) emerge as




Figure 3.4: Stage 3 In-Sample Permutation Feature Importance Rankings across Modalities

the dominant predictive drivers, while sentiment indicators provide moderate secondary refinement, aligning fully with downstream SHAP values (Figure 4.2) and modality ablation evidence (Table 4.3).

3.4 Baseline Econometric & Machine Learning Formulations

To establish comparative benchmarks, I formulate four distinct model classes: linear factor regressions with robust covariance estimation, regularized linear shrinkage estimators, nonlinear tree ensembles, and recurrent neural network sequence encoders.

3.4.1 Fama-MacBeth 2-Pass Linear Regression & Newey-West HAC Estimation

The benchmark linear asset pricing model follows the two-pass cross-sectional regression procedure of Fama and MacBeth (1973), incorporating 1-day lagged factor loadings to prevent lookahead bias and heteroskedasticity and autocorrelation consistent (HAC) covariance matrices (Newey and West, 1987).

Pass 1: Rolling Time-Series Factor Loading Estimation

For each individual stock $i \in \{1, \dots, N_t\}$ on trading day t , factor loadings (betas) are estimated over a rolling historical lookback window of $W = 252$ trading days (one calendar year). Let $\mathbf{R}_{i,t} = [R_{i,t-W+1}, R_{i,t-W+2}, \dots, R_{i,t}]^T \in \mathbb{R}^{W \times 1}$ denote the vector of historical daily excess returns for asset i , and let $\mathbf{F}_t \in \mathbb{R}^{W \times (K+1)}$ represent the matrix of K benchmark factor returns plus a column of ones for the intercept:

$$\mathbf{F}_t = \begin{bmatrix} 1 & F_{1,t-W+1} & F_{2,t-W+1} & \cdots & F_{K,t-W+1} \\ 1 & F_{1,t-W+2} & F_{2,t-W+2} & \cdots & F_{K,t-W+2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & F_{1,t} & F_{2,t} & \cdots & F_{K,t} \end{bmatrix} \quad (3.3)$$

The first-pass time series Ordinary Least Squares (OLS) regression for stock i is formulated as:

$$\mathbf{R}_{i,t} = \mathbf{F}_t \boldsymbol{\beta}_{i,t} + \boldsymbol{\epsilon}_{i,t} \quad (3.4)$$

where $\boldsymbol{\beta}_{i,t} = [\alpha_{i,t}, \beta_{i,1,t}, \beta_{i,2,t}, \dots, \beta_{i,K,t}]^T \in \mathbb{R}^{(K+1) \times 1}$. The OLS point estimator is obtained by minimizing the sum of squared residuals:

$$\hat{\boldsymbol{\beta}}_{i,t} = \arg \min_{\boldsymbol{\beta}} (\mathbf{R}_{i,t} - \mathbf{F}_t \boldsymbol{\beta})^T (\mathbf{R}_{i,t} - \mathbf{F}_t \boldsymbol{\beta}) = (\mathbf{F}_t^T \mathbf{F}_t)^{-1} \mathbf{F}_t^T \mathbf{R}_{i,t} \quad (3.5)$$

To enforce zero-lookahead protection, factor loadings $\hat{\boldsymbol{\beta}}_{i,t}$ estimated using data up to day t are indexed to predict returns at day $t + 1$.

Pass 2: Daily Cross-Sectional Factor Risk Premia Estimation

On each trading day $t + 1$, a cross-sectional OLS regression is estimated across all active constituent assets N_{t+1} using the 1-day lagged factor loadings $\hat{\mathbf{B}}_t = [\mathbf{1}, \hat{\boldsymbol{\beta}}_{1,t}, \hat{\boldsymbol{\beta}}_{2,t}, \dots, \hat{\boldsymbol{\beta}}_{N_{t+1},t}]^T \in \mathbb{R}^{N_{t+1} \times (K+1)}$:

$$\mathbf{R}_{t+1} = \hat{\mathbf{B}}_t \boldsymbol{\gamma}_{t+1} + \boldsymbol{\eta}_{t+1} \quad (3.6)$$

where $\mathbf{R}_{t+1} = [R_{1,t+1}, R_{2,t+1}, \dots, R_{N_{t+1},t+1}]^T \in \mathbb{R}^{N_{t+1} \times 1}$ is the cross-section of realized forward returns, $\boldsymbol{\gamma}_{t+1} = [\gamma_{0,t+1}, \gamma_{1,t+1}, \dots, \gamma_{K,t+1}]^T \in \mathbb{R}^{(K+1) \times 1}$ represents the vector of daily factor risk premia (and cross-sectional intercept $\gamma_{0,t+1}$), and $\boldsymbol{\eta}_{t+1}$ is the vector of cross-sectional pricing errors.

The closed-form daily cross-sectional risk premium estimator is given by:

$$\hat{\boldsymbol{\gamma}}_{t+1} = \left(\hat{\mathbf{B}}_t^T \hat{\mathbf{B}}_t \right)^{-1} \hat{\mathbf{B}}_t^T \mathbf{R}_{t+1} \quad (3.7)$$

Time-Series Aggregation & Step-by-Step Newey-West HAC Proof

Over an evaluation period of T trading days, the overall factor risk premium vector is estimated as the sample mean across daily cross-sectional estimates:

$$\bar{\boldsymbol{\gamma}} = \frac{1}{T} \sum_{t=1}^T \hat{\boldsymbol{\gamma}}_t \quad (3.8)$$

Because daily factor risk premia estimates $\hat{\boldsymbol{\gamma}}_t$ suffer from serial correlation (autocorrelation) due to overlapping return measurement horizons and conditional heteroskedasticity in asset volatility, standard i.i.d. variance estimators underestimate true standard errors, inflating test statistics.

To derive a consistent estimator for the asymptotic covariance matrix of $\bar{\gamma}$, let $\mathbf{u}_t = \hat{\gamma}_t - \bar{\gamma} \in \mathbb{R}^{(K+1) \times 1}$ denote the vector of mean-zero risk premium innovations. The variance of the sample average factor risk premium is:

$$\text{Var}(\bar{\gamma}) = \text{Var}\left(\frac{1}{T} \sum_{t=1}^T \hat{\gamma}_t\right) = \frac{1}{T^2} \sum_{t=1}^T \sum_{s=1}^T \mathbb{E}[\mathbf{u}_t \mathbf{u}_s^T] \quad (3.9)$$

Defining the sample autocovariance matrix at lag j as:

$$\hat{\mathbf{\Gamma}}_j = \frac{1}{T} \sum_{t=j+1}^T \mathbf{u}_t \mathbf{u}_{t-j}^T, \quad \text{for } j \geq 0 \quad (3.10)$$

Equation (3.9) can be expanded into the sum of zero-lag covariance and off-diagonal cross-lag covariances:

$$\text{Var}(\bar{\gamma}) = \frac{1}{T} \left[\hat{\mathbf{\Gamma}}_0 + \sum_{j=1}^{T-1} \left(\hat{\mathbf{\Gamma}}_j + \hat{\mathbf{\Gamma}}_j^T \right) \right] \quad (3.11)$$

Unweighted truncation of lag covariances at finite lag $L < T$ can lead to non-positive semi-definite covariance matrices in finite samples. The Newey and West (1987) HAC covariance matrix estimator guarantees positive semi-definiteness by weighting lag autocovariances with the triangular Bartlett kernel $w(j, L)$:

$$\hat{\mathbf{\Omega}}_{\text{HAC}} = \hat{\mathbf{\Gamma}}_0 + \sum_{j=1}^L w(j, L) \left(\hat{\mathbf{\Gamma}}_j + \hat{\mathbf{\Gamma}}_j^T \right) \quad (3.12)$$

where the Bartlett kernel weights are defined as:

$$w(j, L) = 1 - \frac{j}{L+1}, \quad j = 1, 2, \dots, L \quad (3.13)$$

and the lag truncation bandwidth L is set according to the bandwidth selection rule $L = \lfloor 4(T/100)^{2/9} \rfloor$.

Proof of Guaranteed Positive Semi-Definiteness. To prove that $\hat{\mathbf{\Omega}}_{\text{HAC}} \succeq 0$, express the Bartlett kernel as the spectral window resulting from the convolution of a rectangular