

**UNIVERSITÀ DEGLI STUDI DI VERONA**

DEPARTMENT OF ECONOMICS

Master's Degree in Economics and Data Analysis

**DO MACHINE LEARNING MODELS IMPROVE  
STOCK RETURN PREDICTION?**

**Evidence from S&P 500 Constituent Markets and Dimension Sensitivity  
(2015–2024)**

**Supervisor:**

Prof. Giuseppina Chesini

**Candidate:**

Murtuza Yusuf Rangwala

Matricola: VR508566

Academic Year 2024/2025

## Abstract

Do complex machine learning algorithms actually outperform classical econometric benchmarks when predicting daily US stock returns? I address this empirical question by testing a broad range of quantitative models across a 10-year point-in-time panel of S&P 500 constituent equities from 2015 through 2024 (2,516 daily cross-sections). The predictor universe comprises 59 characteristics spanning technical indicators, quarterly Bloomberg accounting ratios, macroeconomic state variables, news/social sentiment scores, and rolling factor risk exposures.

To prevent lookahead leakage, model parameters and hyperparameter search spaces are optimized strictly inside a 5-fold expanding-window walk-forward validation protocol across the 2020–2024 out-of-sample period. Empirical results confirm that deep sequence models outperform linear estimators (Fama-MacBeth OLS, ElasticNet, LASSO) and static decision tree ensembles (Random Forest, XGBoost). My custom **Temporal Fusion Deep Multimodal Gated Attention (TFDMGA)** architecture delivers the highest predictive accuracy, recording an out-of-sample daily rank Information Coefficient ( $IC$ ) of +0.0348, an Information Ratio ( $ICIR$ ) of 3.12, a ROC AUC of 0.6120, and a directional accuracy of 56.84%. A Diebold-Mariano test with Harvey-Leybourne-Newbold small-sample corrections verifies that these forecast gains over baseline PyTorch LSTMs are statistically significant ( $DM = 2.41, p = 0.016$ ).

When deployed in simulated daily trading portfolios incorporating 10 bps round-trip transaction frictions and a 1-day execution buffer, an initial \$1,000 USD deposit starting in January 2020 grows to \$3,120.99 USD under a baseline LSTM long strategy. Adding a 2:1 Take-Profit/Stop-Loss risk overlay expands capital accumulation to \$4,368.50 USD for LSTM and \$6,482.10 USD for TFDMGA. However, Fama-French 5-factor spanning regressions yield an estimated net strategy alpha of  $\hat{\alpha} = -0.18\%$  per year ( $p = 0.976, R^2 = 41.2\%$ ). Because net alpha is statistically zero, strategy returns

generated by these deep learning models represent dynamic compensation for systematic factor risk exposures rather than unpriced market inefficiency.

**Keywords:** Empirical Asset Pricing, Cross-Sectional Return Prediction, Deep Learning, Transformer Attention, Fama-French Spanning, Transaction Costs, Walk-Forward Validation.

يا رب زدني علما

*“And say, ‘My Lord, increase me in knowledge.’ ”*

*, The Holy Qur’an, Surah Ta-Ha (20:114)*

## Acknowledgements

I express my sincere gratitude to my supervisor, **Prof. Giuseppina Chesini**. Her mentorship, academic guidance, and critical feedback throughout this research project shaped the empirical design and econometric focus of this Master's thesis.

I extend my thanks to the faculty and academic staff of the **Department of Economics at the University of Verona**. The coursework and quantitative training provided during my Master's studies laid the foundation for this work.

Moving from Mumbai to Verona to pursue my Master's degree has been a deeply rewarding journey. I am thankful to the University of Verona for offering an inspiring academic environment.

I acknowledge the financial support provided by **ESU Verona and the Borsa di Studio scholarship programme**, which enabled me to focus fully on my academic studies.

To my friends and classmates in Italy, thank you for your camaraderie, study sessions, and support during our time together in Verona.

I also recognize the open-source scientific community whose software libraries and computational frameworks made the empirical implementation of this study possible.

Above all, I thank my family. I am forever grateful to my father for his unwavering belief in my potential, his sacrifices, and his constant encouragement to pursue higher education. My gratitude also goes to my mother, my sister, and my extended family for their patience, warmth, and steady support.

**Murtuza Rangwala**

Master of Economics and Data Analysis  
University of Verona

# Contents

|                                                                                     |            |
|-------------------------------------------------------------------------------------|------------|
| <b>Abstract</b>                                                                     | <b>i</b>   |
| <b>Dedication</b>                                                                   | <b>iii</b> |
| <b>Acknowledgements</b>                                                             | <b>iv</b>  |
| <b>1 Introduction</b>                                                               | <b>1</b>   |
| 1.1 Background and Economic Motivation . . . . .                                    | 1          |
| 1.2 Research Questions and Core Hypotheses . . . . .                                | 2          |
| 1.3 Summary of Key Empirical Findings . . . . .                                     | 3          |
| 1.4 Structure of the Thesis Manuscript . . . . .                                    | 4          |
| <b>2 Literature Review</b>                                                          | <b>6</b>   |
| 2.1 Introduction . . . . .                                                          | 6          |
| 2.2 Theoretical Asset Pricing Foundations . . . . .                                 | 7          |
| 2.2.1 Equilibrium Foundations: The CAPM and Arbitrage Pricing<br>Theory . . . . .   | 7          |
| 2.2.2 The Fama-MacBeth Two-Pass Regression Framework . . . . .                      | 8          |
| 2.2.3 Multi-Factor Models: Fama-French, Carhart, and Pastor-<br>Stambaugh . . . . . | 9          |
| 2.3 The Factor Zoo Controversy, Multiple Testing, and Shrinkage . . . . .           | 10         |
| 2.3.1 The Proliferation of Anomalies and Cochrane's "Factor Zoo" . .                | 10         |
| 2.3.2 Multiple Testing Inflation and False Discovery Rates . . . . .                | 11         |

|          |                                                                                     |           |
|----------|-------------------------------------------------------------------------------------|-----------|
| 2.3.3    | High-Dimensional Regularized SDF Estimation . . . . .                               | 11        |
| 2.4      | High-Dimensional SDF Estimation & Characteristic Interactions . . . . .             | 12        |
| 2.4.1    | Non-Parametric Characteristic Pruning via Group LASSO . . . . .                     | 12        |
| 2.4.2    | Sparse Signals and Latent Factor Shrinkage . . . . .                                | 13        |
| 2.5      | Machine Learning and Deep Learning Paradigms . . . . .                              | 14        |
| 2.5.1    | Empirical Asset Pricing Benchmark: Gu, Kelly, and Xiu (2020) . . . . .              | 14        |
| 2.5.2    | Characteristics as Dynamic Risk Exposures: IPCA . . . . .                           | 15        |
| 2.5.3    | Deep Learning and Non-Parametric SDF Estimation . . . . .                           | 15        |
| 2.5.4    | Economic Frictions and Limits to Arbitrage . . . . .                                | 16        |
| 2.6      | Financial Time-Series Econometrics & Inference . . . . .                            | 16        |
| 2.6.1    | Heteroskedasticity and Autocorrelation Consistent (HAC) Infer-<br>ence . . . . .    | 16        |
| 2.6.2    | Out-of-Sample Predictive Accuracy: Diebold-Mariano and HLN<br>Corrections . . . . . | 17        |
| 2.6.3    | Data Mining Diagnostics: Reality Check and Superior Predictive<br>Ability . . . . . | 18        |
| 2.7      | Financial ML Protocols & Overfitting Mitigation . . . . .                           | 18        |
| 2.7.1    | Methodological Pitfalls in Financial Machine Learning . . . . .                     | 18        |
| 2.7.2    | Validation Protocols . . . . .                                                      | 19        |
| 2.7.3    | Microstructure Frictions & Transaction Cost Taxonomy . . . . .                      | 19        |
| 2.8      | Synthesis of Core Literature Benchmarks . . . . .                                   | 19        |
| 2.9      | Summary . . . . .                                                                   | 20        |
| <b>3</b> | <b>Data &amp; Methodology</b>                                                       | <b>21</b> |
| 3.1      | Data Pipeline and Sample Construction . . . . .                                     | 21        |

|          |                                                                                        |           |
|----------|----------------------------------------------------------------------------------------|-----------|
| 3.1.1    | Point-in-Time Bloomberg Filing Date Alignment . . . . .                                | 22        |
| 3.1.2    | Survivorship-Bias-Free Dynamic Panel . . . . .                                         | 22        |
| 3.2      | Feature Space Taxonomy and Selection . . . . .                                         | 23        |
| 3.2.1    | Economic Rationale for Feature Set Size (59 Variables vs. 120+ Zoo Features) . . . . . | 23        |
| 3.3      | Feature Normalization and Selection Protocol . . . . .                                 | 24        |
| 3.4      | Baseline Econometric & Machine Learning Formulations . . . . .                         | 26        |
| 3.4.1    | Fama-MacBeth 2-Pass Linear Regression & Newey-West HAC Estimation . . . . .            | 27        |
| 3.4.2    | Regularized Linear Models: ElasticNet, LASSO, & Ridge . . . . .                        | 31        |
| 3.4.3    | Non-Linear Tree Ensembles: Random Forest & XGBoost . . . . .                           | 34        |
| 3.4.4    | Baseline Sequence Encoder: PyTorch LSTM . . . . .                                      | 38        |
| 3.5      | Custom Architecture: TFDMGA Network . . . . .                                          | 39        |
| 3.5.1    | Temporal Feature Encoding via Causal 1D Dilated TCN . . . . .                          | 40        |
| 3.5.2    | Sequential Directed Ring Attention Cascade . . . . .                                   | 41        |
| 3.5.3    | 3-Way Macro Dynamic Gating Network . . . . .                                           | 42        |
| 3.5.4    | Residual Fusion Block & Pre-LN Transformer Encoder . . . . .                           | 44        |
| 3.5.5    | Multi-Task Multi-Head Objective Function . . . . .                                     | 45        |
| 3.6      | Cross-Validation & Walk-Forward Evaluation Scheme . . . . .                            | 46        |
| 3.7      | Hyperparameter Search Space & Training Setup . . . . .                                 | 47        |
| 3.8      | Computational Complexity and Resource Profile . . . . .                                | 48        |
| <b>4</b> | <b>Empirical Results &amp; Backtesting</b>                                             | <b>49</b> |
| 4.1      | Introduction . . . . .                                                                 | 49        |
| 4.2      | Contemporaneous Fama-MacBeth OLS Baseline . . . . .                                    | 49        |



|          |                                                                                          |           |
|----------|------------------------------------------------------------------------------------------|-----------|
| 4.3      | Mathematical Derivation & Empirical Reconciliation of LASSO Shrinkage Collapse . . . . . | 50        |
| 4.4      | SHAP Feature Interpretability Analysis . . . . .                                         | 51        |
| 4.5      | Out-of-Sample Predictive Evaluation . . . . .                                            | 53        |
| 4.6      | Ablation Study of Neural Architecture Components . . . . .                               | 54        |
| 4.7      | Robustness Checks and Sensitivity Analysis . . . . .                                     | 54        |
| 4.8      | Portfolio Construction & Execution Friction . . . . .                                    | 55        |
| 4.9      | Transaction Cost Sensitivity & \$1,000 USD Account Compounding . . .                     | 55        |
| 4.10     | Fama-French 5-Factor Spanning Regressions . . . . .                                      | 57        |
| <b>5</b> | <b>Conclusion &amp; Future Extensions</b>                                                | <b>59</b> |
| 5.1      | Summary of Research Findings . . . . .                                                   | 59        |
| 5.2      | Theoretical and Practical Implications . . . . .                                         | 60        |
| 5.3      | Limitations of the Study . . . . .                                                       | 61        |
| 5.4      | Future Research Extensions . . . . .                                                     | 61        |
|          | <b>Bibliography</b>                                                                      | <b>63</b> |

## List of Figures

|     |                                                                                                         |    |
|-----|---------------------------------------------------------------------------------------------------------|----|
| 3.1 | Bloomberg Point-in-Time Data Pipeline and Staleness Truncation Flowchart . . . . .                      | 23 |
| 3.2 | Variable Taxonomy and Feature Group Modality Split (53 ML Inputs + 6 Factor Betas = 59 Total) . . . . . | 24 |
| 3.3 | Spearman Rank Correlation Heatmap of Selected 53 ML Features . . . .                                    | 25 |
| 3.4 | Stage 3 In-Sample Permutation Feature Importance Rankings across Modalities . . . . .                   | 26 |
| 3.5 | Architecture Diagram of the Temporal Fusion Deep Multimodal Gated Attention (TFDMGA) Network . . . . .  | 39 |
| 3.6 | 5-Fold Expanding-Window Walk-Forward Validation Protocol (2015–2024) . . . . .                          | 46 |
| 4.1 | Global SHAP Mean Absolute Feature Importance Rankings (21-Day Horizon) . . . . .                        | 51 |
| 4.2 | SHAP Summary Beeswarm Plot Illustrating Feature Value Impact on Model Predictions . . . . .             | 52 |
| 4.3 | Quintile Portfolio Sorting, 1-Day Execution Delay, and Turnover Rebalancing Flowchart . . . . .         | 55 |
| 4.4 | Out-of-Sample Cumulative Return Trajectories (Net 10 bps, 2020–2024 Folds) . . . . .                    | 56 |
| 4.5 | Rolling 126-Day Annualized Sharpe Ratios across Model Architectures                                     | 56 |
| 4.6 | Impact of 2:1 Take-Profit/Stop-Loss Risk Overlay (+4% / -2%) on Portfolio Returns . . . . .             | 56 |

|     |                                                                                          |    |
|-----|------------------------------------------------------------------------------------------|----|
| 4.7 | Maximum Drawdown Mitigation During Macro Stress Periods (COVID<br>2020 & Fed Hikes 2022) | 57 |
|-----|------------------------------------------------------------------------------------------|----|

## List of Tables

|     |                                                                                                             |    |
|-----|-------------------------------------------------------------------------------------------------------------|----|
| 2.1 | Comprehensive Synthesis of Core Literature Benchmarks . . . . .                                             | 20 |
| 3.1 | Hyperparameter Search Space and Optimal Tuned Parameters Across<br>Models . . . . .                         | 47 |
| 3.2 | Computational Complexity, Parameter Count, and Resource Execution<br>Profile . . . . .                      | 48 |
| 4.1 | Fama-MacBeth Daily Cross-Sectional Regression Risk Premia (2015–<br>2024) . . . . .                         | 49 |
| 4.2 | Out-of-Sample Predictive Performance Metrics (2020–2024 Test Folds –<br>Master Panel 53 Features) . . . . . | 53 |
| 4.3 | Systematic Component Ablation Study of the TFDMGA Network . . . .                                           | 54 |
| 4.4 | Predictive Accuracy and Account Compounding Across Prediction Hori-<br>zons and Fees . . . . .              | 55 |
| 4.5 | \$1,000 USD Account Growth Under Transaction Cost Scenarios (2020–<br>2024) . . . . .                       | 55 |
| 4.6 | Fama-French 5-Factor Spanning Regressions with Newey-West HAC<br>Errors . . . . .                           | 57 |

# 1. Introduction

## 1.1 Background and Economic Motivation

Why do financial asset returns vary across stocks, and how effectively do market prices incorporate available information? Answering this fundamental question lies at the heart of empirical asset pricing (Fama, 1970; Fama and French, 1993). For decades, academic researchers modeled expected equity returns through linear factor frameworks, evolving from Markowitz’s (1952) mean-variance framework to the Capital Asset Pricing Model (CAPM) of Sharpe (1964) and Lintner (1965), Ross’s (1976) Arbitrage Pricing Theory, Carhart’s (1997) four-factor momentum specification, Pástor and Stambaugh’s (2003) liquidity model, Hou, Xue, and Zhang’s (2015)  $q$ -factor framework, and Fama and French’s (2015) five-factor model. These classic models share a core assumption: stock returns are governed linearly by a small set of static risk exposures.

Real-world financial markets, however, display non-linearities, heavy-tailed return distributions, and dynamic regime shifts. Firm-level characteristics interact in nonlinear ways, for instance, valuation metrics behave differently in market drawdowns compared to economic expansions (Gu et al., 2020; Kozak et al., 2020). Standard Ordinary Least Squares (OLS) regressions face structural hurdles in high dimensional panels, risking omitted variable bias when specifications are too simple, or variance inflation when dozens of firm characteristics are included simultaneously.

Machine learning (ML) and deep sequence architectures provide a flexible framework for modeling nonlinear characteristic interactions and temporal patterns in equity markets (Gu et al., 2020, 2021; Chen et al., 2024; Kozak et al., 2020; Freyberger et al., 2020; Kelly et al., 2019). Techniques ranging from regularized linear models (Tibshirani, 1996; Zou and Hastie, 2005) and decision tree ensembles (Breiman, 2001; Friedman, 2001; Chen and Guestrin, 2016) to temporal convolutional networks (Bai et al., 2018) and multimodal Transformer architectures (Vaswani et al., 2017; Lim et al., 2021) allow researchers to extract predictive signals directly from raw panel data.

Yet applying machine learning to financial asset pricing differs from standard computer vision or language modeling. Financial return series are characterized by low signal-to-noise ratios, regime shifts, and structural breaks (López de Prado, 2018, 2020). Naive applications risk lookahead bias, survivorship bias, multiple testing inflation (Harvey et al., 2016; Harvey and Liu, 2020; Hou et al., 2020), and unmodeled transaction cost drag (Novy-Marx and Velikov, 2016; McLean and Pontiff, 2016; Avramov et al., 2023). Without out-of-sample walk-forward validation and trading friction modeling, apparent predictive gains can quickly disappear in live trading.

In this thesis, I evaluate machine learning models for cross-sectional return prediction across S&P 500 equities over the 2015–2024 period, addressing these methodological challenges directly.

## 1.2 Research Questions and Core Hypotheses

To evaluate the statistical predictability and economic viability of machine learning models, I frame my investigation around four central research questions:

**Research Question 1 (Statistical Predictability):** Do nonlinear machine learning models (Random Forest, XGBoost, LSTM, TFDMDGA) achieve statistically significant out-of-sample return predictability compared to linear econometric baselines (Fama-MacBeth OLS, LASSO, ElasticNet)?

*Hypothesis 1:* Nonlinear models, especially sequence-aware deep neural networks (Bai et al., 2018; Chen et al., 2024), outperform linear baselines in out-of-sample Information Coefficient ( $IC$ ), Information Ratio ( $ICIR$ ), and Directional Accuracy by capturing nonparametric characteristic interactions and temporal dependencies.

**Research Question 2 (Multi-Modal Information Integration):** Does integrating technical indicators, quarterly Bloomberg accounting ratios, macroeconomic state variables, and news/social sentiment within a dynamic gating architecture improve prediction accuracy over single-modality models?

*Hypothesis 2:* Multimodal feature integration (Cong et al., 2021; Lim et al., 2021) allows dynamic attention mechanisms to reweight feature modalities as macroeconomic regimes shift, yielding higher out-of-sample correlation and signal stability.

**Research Question 3 (Economic Significance & Transaction Friction):** Can out-of-sample return forecasts generate profitable trading strategies after accounting for transaction frictions (10 bps round-trip fees, execution delays, turnover, and stop-loss risk management)?

*Hypothesis 3:* While daily portfolio rebalancing incurs cost drag under 10 bps transaction fees (Novy-Marx and Velikov, 2016; Avramov et al., 2023), deep sequence models paired with dynamic take-profit/stop-loss overlays help control drawdowns and compound capital over time.

**Research Question 4 (Econometric Factor Spanning & Market Efficiency):** Are out-of-sample returns produced by deep learning models indicative of unpriced market alpha ( $\alpha > 0$ ), or are they fully spanned by systematic risk factors ( $\alpha = 0$ ) within the Fama-French 5-Factor framework?

*Hypothesis 4:* Out-of-sample strategy returns are explained by systematic factor risk exposures ( $Mkt - RF, SMB, HML, RMW, CMA$ ) (Fama and French, 2015; Kelly et al., 2019), indicating that machine learning models act as dynamic factor allocation engines rather than discoverers of unpriced arbitrage.

### 1.3 Summary of Key Empirical Findings

Using a 5-fold expanding-window walk-forward validation scheme over the 2020–2024 out-of-sample period (detailed in Chapter 3 and Chapter 4), my main findings are:

1. **Performance in Return Rank Forecasting:** The custom Temporal Fusion Deep Multimodal Gated Attention (**TFDMGA**) network achieves the highest predictive

accuracy among all evaluated models, recording an out-of-sample daily Information Coefficient of  $IC = +0.0348$ , an Information Ratio of  $ICIR = 3.12$ , a ROC AUC of 0.6120, and a directional accuracy of 56.84%.

2. **Statistical Comparison with Baseline Sequence Models:** A formal Diebold and Mariano (1995) forecast accuracy test with Harvey et al. (1997) small-sample corrections confirms that TFDMGA yields statistically significant performance gains over a baseline PyTorch LSTM model ( $DM = 2.41, p = 0.016$ ), validating the structural benefits of crossmodal ring attention and dynamic macro gating.
3. **Capital Compounding Trajectory:** Under a realistic 10 bps round-trip transaction cost setting and a 1-day execution buffer (Novy-Marx and Velikov, 2016), an initial \$1,000 USD deposit in January 2020 grows to **\$3,120.99 USD** under the baseline LSTM  $Q_5$  long strategy, and reaches **\$4,368.50 USD** with a 2:1 Take-Profit/Stop-Loss (TPSL) risk overlay. Under TFDMGA with the TPSL overlay, capital compounds to **\$6,482.10 USD** by December 2024 (see Table 4.5 in Chapter 4).
4. **Fama-French 5-Factor Spanning & Factor Timing:** Spanning regressions against the Fama and French (2015) 5-factor model yield an estimated net strategy alpha of  $\hat{\alpha} = -0.18\%$  per year ( $p = 0.976, R^2 = 41.2\%$ ). Because net alpha is statistically zero, strategy returns are explained by dynamic risk exposure to systematic factor risks, specifically operating profitability ( $RMW$  beta  $= +0.512, t = +9.12$ ), consistent with market efficiency (Fama, 1970).

## 1.4 Structure of the Thesis Manuscript

The remainder of this thesis is structured across four consolidated chapters:

- **Chapter 2 (Literature Review)** surveys academic literature in empirical asset pricing, factor models, high dimensional regularization, and financial machine learning.



- **Chapter 3 (Data & Methodology)** details the Bloomberg point-in-time data pipeline, constituent panel construction, 59-variable feature taxonomy, baseline model formulations, and the custom TFDMGA network architecture.
- **Chapter 4 (Empirical Results & Backtesting)** presents daily Fama-MacBeth baselines, LASSO shrinkage derivations, SHAP interpretability analysis, out-of-sample predictive metrics, ablation results, trading backtests under transaction fees, and Fama-French spanning regressions.
- **Chapter 5 (Conclusion & Future Research Extensions)** summarizes the main findings, discusses theoretical and practical implications, addresses study limitations, and outlines future research extensions.

## 2. Literature Review

### 2.1 Introduction

Understanding stock return predictability requires tracing the evolution of empirical asset pricing over the past six decades. Early academic frameworks established that asset prices in competitive markets reflect available information (Fama, 1970), modeling expected excess returns as linear functions of systematic risk exposures (Sharpe, 1964; Lintner, 1965). Subsequent econometric advances operationalized these principles through multifactor cross-sectional regression models (Fama and MacBeth, 1973; Ross, 1976; Fama and French, 1993, 2015).

However, the rapid expansion of financial data led to what Cochrane (2011) famously termed the “factor zoo”, a proliferation of over 300 candidate anomaly variables claiming to generate alpha. In response, modern financial econometrics has developed along two primary paths: (i) statistical auditing of data-mining inflation and false discovery rates (Harvey et al., 2016; Harvey and Liu, 2020; Hou et al., 2020), and (ii) high dimensional regularized estimation and nonparametric machine learning (Kozak et al., 2020; Freyberger et al., 2020; Gu et al., 2020; Chen et al., 2024).

In this chapter, I synthesize the asset pricing literature around six key thematic pillars that ground the theoretical framework, model architectures, feature selection, and statistical inference used in this thesis:

1. **Theoretical Asset Pricing Foundations:** From the single factor equilibrium of CAPM and APT to the multifactor models of Fama-French, Carhart, and Pastor-Stambaugh.
2. **The Factor Zoo, Multiple Testing, and Shrinkage:** Multiple testing inflation, false discovery rates, and high dimensional regularized SDF estimation.

3. **High-Dimensional SDF Estimation & Characteristic Interactions:** Nonparametric characteristic selection, latent factor models, and dynamic risk loadings.
4. **Machine Learning & Deep Learning Paradigms:** Nonlinear empirical asset pricing, instrumented principal components, and deep sequence SDF models.
5. **Financial Time-Series Econometrics & Inference:** HAC covariance estimation, out-of-sample forecast accuracy tests, and data-mining diagnostics.
6. **Financial ML Protocols & Overfitting Mitigation:** Backtesting rules, expanding walk-forward splits, microstructural frictions, and transaction cost frameworks.

## 2.2 Theoretical Asset Pricing Foundations

### 2.2.1 Equilibrium Foundations: The CAPM and Arbitrage Pricing Theory

Modern quantitative finance begins with Markowitz’s (1952) mean-variance framework and the Capital Asset Pricing Model (CAPM) of Sharpe (1964) and Lintner (1965). CAPM assumes homogeneous investor expectations, frictionless markets, and quadratic utility. Under these ideal conditions, every investor holds the market portfolio  $R_m$ , making it mean-variance efficient. An asset’s expected excess return is governed strictly by its market beta:

$$\mathbb{E}[R_{i,t}] - R_{f,t} = \beta_{i,m} (\mathbb{E}[R_{m,t}] - R_{f,t}) \quad (2.1)$$

where  $\beta_{i,m} = \frac{\text{Cov}(R_{i,t}, R_{m,t})}{\text{Var}(R_{m,t})}$  captures systemic market risk. Black (1972) relaxed risk-free borrowing assumptions, introducing the zero-beta CAPM. Black showed that borrowing constraints flatten the Security Market Line while raising its intercept.

To address CAPM’s equilibrium restrictions, Ross (1976) proposed Arbitrage Pricing Theory (APT). Grounded in asymptotic arbitrage-free pricing rather than market equilibrium, APT specifies a linear  $K$ -factor generative return structure:

$$R_{i,t} = \mathbb{E}[R_{i,t}] + \sum_{k=1}^K \beta_{i,k} f_{k,t} + \epsilon_{i,t} \quad (2.2)$$

where  $f_{k,t}$  represent mean-zero common macroeconomic risk factors,  $\beta_{i,k}$  are factor loadings, and  $\epsilon_{i,t}$  is idiosyncratic noise with  $\text{Cov}(\epsilon_{i,t}, \epsilon_{j,t}) = 0$  for  $i \neq j$ . Applying linear pricing space arguments in Hilbert spaces, Ross established that expected returns are linearly spanned by factor risk premia  $\lambda_k$ :

$$\mathbb{E}[R_{i,t}] - R_{f,t} = \sum_{k=1}^K \beta_{i,k} \lambda_k \quad (2.3)$$

### 2.2.2 The Fama-MacBeth Two-Pass Regression Framework

To test whether factor loadings explain cross-sectional stock returns, Fama and MacBeth (1973) introduced a two-pass regression methodology. In Pass 1, I run time-series regressions asset by asset over a rolling historical lookback window  $T_1$  to estimate factor betas  $\hat{\beta}_{i,k}$ :

$$R_{i,t} - R_{f,t} = \alpha_i + \sum_{k=1}^K \beta_{i,k} f_{k,t} + \epsilon_{i,t}, \quad t = 1, \dots, T_1 \quad (2.4)$$

In Pass 2, I run daily cross-sectional regressions across all available equities  $i = 1, \dots, N$  using estimated factor betas as explanatory variables:

$$R_{i,t} - R_{f,t} = \gamma_{0,t} + \sum_{k=1}^K \hat{\beta}_{i,k} \gamma_{k,t} + u_{i,t}, \quad i = 1, \dots, N \quad (2.5)$$

The point estimates of factor risk premia  $\bar{\gamma}_k$  and their standard errors are obtained by averaging cross-sectional slope coefficients over time:

$$\bar{\gamma}_k = \frac{1}{T_2} \sum_{t=1}^{T_2} \hat{\gamma}_{k,t}, \quad \text{SE}(\bar{\gamma}_k) = \sqrt{\frac{1}{T_2(T_2 - 1)} \sum_{t=1}^{T_2} (\hat{\gamma}_{k,t} - \bar{\gamma}_k)^2} \quad (2.6)$$

Because standard Fama-MacBeth standard errors ignore serial correlation and generated-regressor error, modern applications apply Heteroskedasticity and Autocorrelation Consistent (HAC) corrections (Newey and West, 1987) and Shanken (1992) beta-uncertainty adjustments.

### 2.2.3 Multi-Factor Models: Fama-French, Carhart, and Pastor-Stambaugh

Single-factor CAPM faced empirical challenges in the 1980s as market anomalies accumulated. Fama and French (1992) documented that size (market capitalization) and book-to-market ( $B/M$ ) ratios absorbed much of market beta's explanatory power. To capture these anomalies, Fama and French (1993) formulated the Fama-French Three-Factor Model:

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2}SMB_t + \beta_{i,3}HML_t + e_{i,t} \quad (2.7)$$

where  $SMB_t$  (Small Minus Big) isolates the size premium, and  $HML_t$  (High Minus Low) isolates the value premium.

To address the price momentum anomaly (Jegadeesh and Titman, 1993), Carhart (1997) added a momentum factor ( $WML_t$ , Winners Minus Losers):

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2}SMB_t + \beta_{i,3}HML_t + \beta_{i,4}WML_t + e_{i,t} \quad (2.8)$$

Concurrently, Pástor and Stambaugh (2003) demonstrated that aggregate market liquidity represents an un-diversifiable risk factor. Stocks with high liquidity sensitivity command higher expected returns to compensate for illiquidity during market drawdowns.

To incorporate operating profitability and investment intensity, Fama and French (2015) introduced their Five-Factor Asset Pricing Model:

$$\begin{aligned} R_{i,t} - R_{f,t} = & \alpha_i + \beta_{i,1}(R_{m,t} - R_{f,t}) + \beta_{i,2}SMB_t + \beta_{i,3}HML_t \\ & + \beta_{i,4}RMW_t + \beta_{i,5}CMA_t + e_{i,t} \end{aligned} \quad (2.9)$$

where  $RMW_t$  (Robust Minus Weak) measures profitability differentials, and  $CMA_t$  (Conservative Minus Aggressive) measures investment intensity. Grounded in the dividend discount model, Fama and French (2015) showed that holding book-to-market constant, higher profitability implies higher expected returns, while higher investment lowers expected returns. In parallel, Hou et al. (2015) introduced the  $q$ -factor model

( $Mkt$ ,  $ME$ ,  $I/A$ ,  $ROE$ ), establishing that investment and profitability factors derived from  $q$ -theory span a broad set of anomalies.

## **2.3 The Factor Zoo Controversy, Multiple Testing, and Shrinkage**

### **2.3.1 The Proliferation of Anomalies and Cochrane's "Factor Zoo"**

By 2011, academic papers claimed over 300 firm characteristics predicted stock returns. In his AFA presidential address, Cochrane (2011) dubbed this proliferation the “factor zoo.” Cochrane noted that while theoretical models assume a small set of macro factors span returns, empirical literature created dozens of redundant predictors without verifying if candidate factors were incremental to existing benchmarks.

This problem was worsened by post-publication return decay. McLean and Pontiff (2016) examined 97 published anomalies and showed that out-of-sample portfolio returns fall by 26% after publication and up to 58% post-sample. This decay reflects arbitrage capital exploiting published signals alongside data-mining bias in original studies.

### 2.3.2 Multiple Testing Inflation and False Discovery Rates

Harvey et al. (2016) analyzed the statistical cause of factor zoo inflation. Standard empirical studies tested candidate factors using single-hypothesis  $t$ -tests ( $|t| > 1.96$ ,  $\alpha = 0.05$ ). Evaluating hundreds of candidate signals on identical datasets creates high false discovery rates.

Harvey et al. (2016) evaluated 300+ anomaly factors under multiple-testing adjustments (Bonferroni, Holm-Bonferroni). They established that keeping the False Discovery Rate (FDR) below 5% requires raising the minimum  $t$ -statistic threshold to:

$$|t| > 3.0 \quad (p\text{-value} < 0.0027) \quad (2.10)$$

Harvey and Liu (2020) extended this via Bayesian FDR frameworks. In a large replication audit, Hou et al. (2020) re-evaluated 452 cross-sectional anomalies under microcap exclusions and value-weighting. They found that 65% of published anomalies fail at  $|t| > 1.96$ , and over 85% fail under  $|t| > 3.0$ .

### 2.3.3 High-Dimensional Regularized SDF Estimation

To address the factor zoo problem structurally, econometricians re-framed asset pricing as a high dimensional regularized Stochastic Discount Factor (SDF) estimation problem. In a stochastic discount factor framework, an asset pricing model specifies a scalar pricing kernel  $M_{t+1}$  such that for all excess returns  $R_{i,t+1}$ :

$$\mathbb{E}_t [M_{t+1} R_{i,t+1}] = 0 \quad (2.11)$$

When the SDF is parameterized as a linear combination of a high dimensional vector of  $K$  candidate characteristics or factor returns  $\mathbf{z}_{t+1}$ , i.e.,  $M_{t+1} = 1 - \mathbf{b}^T(\mathbf{z}_{t+1} - \mathbb{E}[\mathbf{z}_{t+1}])$ , unregularized sample estimators fail due to multicollinearity and  $K \approx T$  sample noise.

Kozak et al. (2020) showed that shrinking the factor weighting vector  $\mathbf{b}$  via combined  $L_1$  (LASSO) and  $L_2$  (Ridge) regularization yields robust, out-of-sample optimal SDF estimators. The regularized SDF objective function is formulated as:

$$\min_{\mathbf{b}} \left\{ \frac{1}{T} \sum_{t=1}^T (1 - \mathbf{b}^T \mathbf{z}_t)^2 + \gamma_1 \|\mathbf{b}\|_1 + \gamma_2 \|\mathbf{b}\|_2^2 \right\} \quad (2.12)$$

Kozak et al. (2020) showed that while unregularized models tend to overfit high-frequency sample noise, penalized SDF estimation extracts a low-dimensional principal component structure that spans expected returns while shrinking noise coefficients to zero.

## 2.4 High-Dimensional SDF Estimation & Characteristic Interactions

### 2.4.1 Non-Parametric Characteristic Pruning via Group LASSO

Parallel to regularized linear SDFs, Freyberger et al. (2020) developed a nonparametric model to systematically select which firm characteristics provide incremental predictive power for expected returns. Using nonparametric B-spline expansions and Group LASSO penalties, their estimator models expected returns as:

$$\mathbb{E}_t[R_{i,t+1}] = \sum_{j=1}^J g_j(c_{i,j,t}) + \epsilon_{i,t+1} \quad (2.13)$$

where  $c_{i,j,t}$  represents the rank-transformed characteristic  $j$  for firm  $i$  at time  $t$ , and  $g_j(\cdot)$  is a nonlinear continuous function. The Group LASSO penalty penalizes the vector of spline coefficients for each characteristic simultaneously:

$$\min_{\{g_j\}} \left\{ \frac{1}{NT} \sum_{i=1}^N \sum_{t=1}^T \left( R_{i,t+1} - \sum_{j=1}^J g_j(c_{i,j,t}) \right)^2 + \lambda \sum_{j=1}^J \sqrt{\dim(g_j)} \|g_j\|_2 \right\} \quad (2.14)$$

Freyberger et al. (2020) showed that out of more than 50 published characteristics, only a small subset (such as short term momentum, market capitalization, operating



profitability, and asset growth) retain independent predictive power once non-linearities and characteristic redundancies are penalized.

#### **2.4.2 Sparse Signals and Latent Factor Shrinkage**

Chinco et al. (2019) examined sparse return predictability at high frequencies, showing that LASSO regression identifies transient, cross-asset predictive signals that standard time series estimators miss. Grounding model selection in econometric diagnostics, Lewellen et al. (2010) provided a critique of cross-sectional asset pricing tests. They showed that testing multifactor models on strong factor-structure portfolios (such as standard 25 Size-B/M sorted portfolios) produces artificially inflated cross-sectional  $R^2$  values. They prescribed rigorous diagnostic standards: expanding test portfolios beyond Size-B/M, forcing theoretical factor risk premia to equal sample factor means, and reporting confidence intervals for cross-sectional  $R^2$ .

Expanding structural factor selection, Feng et al. (2020) and Giglio and Xiu (2021) developed double-testing econometric frameworks combining LASSO factor selection with two-pass regressions. They showed that high dimensional regularization enables researchers to evaluate whether a newly proposed factor adds incremental explanatory power to the existing universe of 300+ factors without suffering from omitted variable bias or data mining contamination.

## 2.5 Machine Learning and Deep Learning Paradigms

### 2.5.1 Empirical Asset Pricing Benchmark: Gu, Kelly, and Xiu (2020)

Gu et al. (2020) established the primary empirical benchmark for financial machine learning. Evaluating a 60-year panel of over 30,000 U.S. equities (1957–2016) with 94 firm characteristics and 8 macroeconomic indicators, Gu, Kelly, and Xiu tested linear OLS, LASSO, ElasticNet, PCR, PLS, Random Forests, Gradient Boosted Trees (GBDT), and Feedforward Neural Networks (MLP).

The non-linear forecasting panel model takes the form:

$$R_{i,t+1} = g^*(\mathbf{x}_{i,t}) + \epsilon_{i,t+1}, \quad \mathbf{x}_{i,t} = \mathbf{c}_{i,t} \otimes \mathbf{d}_t \quad (2.15)$$

where  $\mathbf{c}_{i,t}$  is a  $p$ -dimensional vector of firm characteristics,  $\mathbf{d}_t$  is an  $m$ -dimensional vector of macro variables, and  $\otimes$  denotes the Kronecker product creating interactive feature conditioning.

Gu et al. (2020) established three key empirical results:

1. **Nonlinear Model Advantage:** Tree ensembles and neural networks outperform linear specifications in out-of-sample  $R_{\text{oos}}^2$ , effectively doubling long-short portfolio Sharpe ratios.
2. **Characteristic Interaction Gains:** Most predictive gain stems from non-linear interaction terms (such as size interacted with momentum and liquidity).
3. **Optimal Network Depth:** Shallow networks (1–3 hidden layers) outperform deeper architectures, as deeper networks overfit on low signal-to-noise equity return data.

### 2.5.2 Characteristics as Dynamic Risk Exposures: IPCA

Kelly et al. (2019) provided theoretical backing for why machine learning forecasts succeed. Introducing Instrumented Principal Component Analysis (IPCA), Kelly, Pruitt, and Su modeled asset returns as:

$$R_{i,t+1} = \mathbf{z}_{i,t}^T \Gamma_{\beta} \mathbf{f}_{t+1} + \epsilon_{i,t+1} \quad (2.16)$$

where  $\mathbf{z}_{i,t}$  is a vector of observable firm characteristics,  $\Gamma_{\beta}$  is a mapping matrix connecting characteristics to dynamic factor loadings ( $\beta_{i,t} = \Gamma_{\beta}^T \mathbf{z}_{i,t}$ ), and  $\mathbf{f}_{t+1}$  represents unobservable latent factor returns.

Kelly et al. (2019) showed that firm characteristics proxy for time-varying factor risk exposures. Machine learning return predictors  $R_{i,t+1} = g(\mathbf{z}_{i,t})$  are thus implicitly estimating dynamic latent factor models with time-varying risk loadings.

### 2.5.3 Deep Learning and Non-Parametric SDF Estimation

Chen et al. (2024) extended machine learning to deep architectures, estimating non-parametric asset pricing SDFs directly from high-dimensional equity panels. Parameterizing factor loadings  $\beta(\mathbf{c}_{i,t})$  via Deep Neural Networks (DNNs) combined with Generative Adversarial Networks (GANs), Chen, Pelger, and Zhu showed that deep learning captures macro regime shifts and structural non-linearities missed by linear factor models.

Similarly, Cong et al. (2021) built multimodal deep learning models to process accounting signals alongside high-dimensional financial news text. To interpret deep neural network predictions, Lundberg and Lee (2017) introduced SHapley Additive exPlanations (SHAP), based on cooperative game theory:

$$f(\mathbf{x}) = \phi_0 + \sum_{j=1}^M \phi_j \mathbf{x}_j \quad (2.17)$$

allowing transparent feature attribution in non-linear models.

#### 2.5.4 Economic Frictions and Limits to Arbitrage

Avramov et al. (2023) examined whether machine learning forecast gains survive real-world market frictions. Avramov, Cheng, and Metzker demonstrated that machine learning predictability is concentrated in illiquid, microcap stocks with high distress risk and short-sale constraints. Excluding microcaps or deducting realistic execution fees significantly reduces out-of-sample strategy profits. This highlights the necessity of incorporating transaction cost frictions into financial machine learning evaluation.

### 2.6 Financial Time-Series Econometrics & Inference

#### 2.6.1 Heteroskedasticity and Autocorrelation Consistent (HAC) Inference

Financial return series exhibit volatility clustering, heavy tails, and serial correlation. Standard OLS covariance matrices underestimate parameter variances, producing inflated  $t$ -statistics. To ensure valid statistical inference, Newey and West (1987) developed the Heteroskedasticity and Autocorrelation Consistent (HAC) covariance matrix estimator.

For regression residuals  $e_t$  and regressors  $\mathbf{X}_t$ , defining  $\mathbf{u}_t = \mathbf{X}_t e_t$ , the Newey-West sample covariance matrix  $\hat{\Omega}_{\text{NW}}$  is:

$$\hat{\Omega}_{\text{NW}} = \hat{\Gamma}_0 + \sum_{j=1}^L w(j, L) \left( \hat{\Gamma}_j + \hat{\Gamma}_j^T \right) \quad (2.18)$$

where  $\hat{\Gamma}_j = \frac{1}{T} \sum_{t=j+1}^T \mathbf{u}_t \mathbf{u}_{t-j}^T$  is the sample autocovariance at lag  $j$ , and  $w(j, L) = 1 - \frac{j}{L+1}$  is the Bartlett kernel weighting function over lag length  $L$ . The Bartlett kernel guarantees that  $\hat{\Omega}_{\text{NW}}$  remains positive semi-definite in finite samples.

### 2.6.2 Out-of-Sample Predictive Accuracy: Diebold-Mariano and HLN Corrections

When comparing out-of-sample forecast accuracy between competing models (e.g. TFDMDGA versus PyTorch LSTM), standard independent sample  $t$ -tests fail due to forecast error serial dependence. Diebold and Mariano (1995) introduced a distribution-free statistic for testing equal predictive accuracy.

Let  $e_{1,t}$  and  $e_{2,t}$  denote forecast errors from Model 1 and Model 2 at time  $t$ . The loss differential sequence under Mean Squared Error is:

$$d_t = L(e_{1,t}) - L(e_{2,t}) \quad (2.19)$$

Under the null hypothesis  $H_0 : \mathbb{E}[d_t] = 0$ , the Diebold-Mariano statistic is:

$$DM = \frac{\bar{d}}{\sqrt{\hat{V}(\bar{d})/T}} \xrightarrow{d} \mathcal{N}(0, 1) \quad (2.20)$$

where  $\bar{d} = \frac{1}{T} \sum_{t=1}^T d_t$ , and  $\hat{V}(\bar{d})$  is the Newey-West HAC-adjusted asymptotic variance.

To correct for small-sample size distortion, Harvey et al. (1997) introduced the Harvey-Leybourne-Newbold (HLN) correction. For forecast horizon  $h$ , the HLN-adjusted statistic is:

$$DM^* = DM \times \left[ \frac{T + 1 - 2h + T^{-1}h(h-1)}{T} \right]^{1/2} \quad (2.21)$$

Critical values for  $DM^*$  are evaluated against the Student's  $t$ -distribution with  $(T - 1)$  degrees of freedom.

### 2.6.3 Data Mining Diagnostics: Reality Check and Superior Predictive Ability

When tuning multiple model architectures across hyperparameters, researchers risk selecting a model that performs well purely by chance. To control for data-mining bias, White (2000) developed the Reality Check for Data Mining using a bootstrap procedure over candidate model forecasts.

Hansen (2005) refined this framework into the Superior Predictive Ability (SPA) test, introducing a studentized test statistic with sample-dependent variance adjustments:

$$T_k^{\text{SPA}} = \max_{k=1,\dots,K} \frac{\sqrt{T} \bar{d}_k}{\hat{\omega}_k} \quad (2.22)$$

where  $\hat{\omega}_k^2$  estimates the asymptotic variance of  $\sqrt{T} \bar{d}_k$ . Hansen showed that SPA provides higher power and robustness against uninformative candidate models.

## 2.7 Financial ML Protocols & Overfitting Mitigation

### 2.7.1 Methodological Pitfalls in Financial Machine Learning

In financial machine learning, standard statistical cross-validation techniques (such as standard randomized  $K$ -fold CV) can fail. As detailed by López de Prado (2018), financial time series data violate the independent and identically distributed (i.i.d.) assumption. Applying naive cross-validation induces several methodological problems:

1. **Lookahead Leakage:** Using future information (such as future target returns or lookahead normalized features) during training.
2. **Serial Memory Contamination:** Overlapping target return windows (such as multi-day cumulative returns) allow training samples to overlap with test samples, leaking labels across folds.
3. **Backtest Overfitting:** Iteratively tuning model parameters directly on test set Sharpe ratios until an illusion of profitability is produced.

### 2.7.2 Validation Protocols

To ensure backtesting integrity, López de Prado (2018) established validation protocols:

- **Expanding Walk-Forward Validation:** Models are trained on historical data  $t \in [1, T_{\text{train}}]$  and evaluated on subsequent out-of-sample periods  $t \in [T_{\text{train}} + 1, T_{\text{test}}]$ . The training window expands iteratively over time.
- **Combinatorial Purged Cross-Validation (CPCV):** When non-sequential cross-validation is conducted, training folds adjacent to test folds are *purged* by removing samples whose label windows overlap with test windows. In addition, an *embargo* period is appended after test sets to eliminate serial correlation leakage.
- **Point-in-Time Data Alignment:** Financial fundamentals and accounting metrics are indexed by their actual public SEC disclosure dates (such as 10-K/10-Q filing dates) rather than fiscal period end dates.

### 2.7.3 Microstructure Frictions & Transaction Cost Taxonomy

High-frequency predictive signals can decay rapidly once realistic transaction costs and market microstructure frictions are accounted for (Novy-Marx and Velikov, 2016; Frazzini and Pedersen, 2014). Novy-Marx and Velikov (2016) show that many published anomaly strategies disappear after deducting effective bid-ask spreads, market impact costs, and turnover frictions.

Second, execution realism mandates an execution delay: trading signals generated at market close on day  $t$  cannot be executed until day  $t + 1$  close or open. In addition, leverage constraints and margin costs (Frazzini and Pedersen, 2014) restrict short-selling capacity on distressed equities, reinforcing the necessity of evaluating long-only and cost-adjusted long-short performance.

## 2.8 Synthesis of Core Literature Benchmarks

Table 2.1 provides a synthesis of the literature that anchors the theoretical framing, model architecture, feature engineering, econometric estimation, and validation of this thesis.

Table 2.1: Comprehensive Synthesis of Core Literature Benchmarks

| Thematic Domain                  | Key Literature Citation                                                                                                                                                  | Core Theoretical / Econometric Finding                                                                                                                                                                                                                                                                                                        | Implementation in Thesis                                                                                                                                                              |
|----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Asset Pricing Foundations        | Sharpe (1964), Lintner (1965)<br>Fama and MacBeth (1973)<br>Ross (1976)<br>Fama and French (1993), Fama and French (2015)<br>Carhart (1997), Pástor and Stambaugh (2003) | Single factor CAPM equilibrium; market beta linear pricing<br>Two-pass cross-sectional regression methodology<br>Arbitrage Pricing Theory (APT); multifactor space<br>3-Factor and 5-Factor models ( <i>Mkt</i> , <i>SMB</i> , <i>HML</i> , <i>RMW</i> , <i>CMA</i> )<br>Price momentum ( <i>WML</i> ) and aggregate market liquidity factors | Baseline risk benchmark<br>Baseline linear factor risk premia<br>Linear multifactor baseline<br>Econometric spanning regressions<br>Extended multifactor benchmark                    |
| Factor Zoo & Shrinkage           | Cochrane (2011)<br>Harvey et al. (2016), Harvey and Liu (2020)<br>Hou et al. (2020), McLean and Pontiff (2016)<br>Kozak et al. (2020)<br>Freyberger et al. (2020)        | Factor zoo diagnosis; lack of unified factor structure<br>Multiple testing bias; minimum hurdle rate $ t  > 3.0$<br>65%+ anomaly replication failure; post-publication decay<br>Regularized $L_1/L_2$ SDF shrinkage in high dimensions<br>Nonparametric Group LASSO characteristic pruning                                                    | High dimensional model motivation<br>Feature auditing threshold<br>Microcap exclusions & value weighting<br>Penalized Ridge/ElasticNet baselines<br>Rank-normalized feature selection |
| Machine Learning & Deep Learning | Gu et al. (2020)<br>Kelly et al. (2019)<br>Chen et al. (2024)<br>Avramov et al. (2023)<br>Lundberg and Lee (2017)                                                        | Nonlinear ML performance; interaction terms; 1–3 layer networks<br>Instrumented PCA (IPCA); characteristics as dynamic betas<br>Deep learning nonparametric SDF estimation via GANs<br>ML gains concentrated in illiquid hard-to-arbitrage stocks<br>SHAP additive feature attribution for nonlinear models                                   | Experimental lineup design<br>Latent dynamic factor framing<br>Sequence model architecture<br>Size-stratified empirical evaluation<br>Model interpretability framework                |
| Time-Series Econometrics         | Newey and West (1987)<br>Diebold and Mariano (1995), Harvey et al. (1997)<br>White (2000), Hansen (2005)                                                                 | HAC robust standard errors; Bartlett kernel weighting<br>Out-of-sample forecast accuracy test & HLN correction<br>Reality Check and Superior Predictive Ability (SPA) tests                                                                                                                                                                   | Fama-MacBeth HAC t-statistics<br>Model forecast comparison tests<br>Hyperparameter data-mining audit                                                                                  |
| Financial ML Validation          | López de Prado (2018)<br>Novy-Marx and Velikov (2016)<br>Frazzini and Pedersen (2014)                                                                                    | Financial ML protocols; purging, embargoing, walk-forward<br>Transaction cost taxonomy; effective spreads and impact<br>Betting against beta; leverage constraints and margin drag                                                                                                                                                            | 5-Fold expanding walk-forward split<br>10 bps friction drag evaluation<br>Execution delay & shorting rules                                                                            |

## 2.9 Summary

In conclusion, the evolution of empirical asset pricing shows a clear transition from simple linear factor models toward regularized high dimensional estimation and nonparametric deep sequence architectures. By synthesizing the theoretical insights of factor pricing, the statistical rigor of multiple testing adjustments, the predictive capability of deep sequence networks, and empirical validation protocols, this thesis builds a framework for cross-sectional stock return prediction.



### 3. Data & Methodology

#### 3.1 Data Pipeline and Sample Construction

My empirical sample comprises daily stock prices, trading volumes, and quarterly financial accounting metrics for active and historical S&P 500 constituent equities from December 31, 2014, through December 31, 2024. The active evaluation window spans January 1, 2015, to December 31, 2024, yielding 2,516 trading days and over 1.25 million stock-day observations (Gu et al., 2020).

Daily closing prices, dividend payments, stock splits, and trading volumes are ingested from Yahoo Finance and cross-verified against Bloomberg Terminal records. Daily risk-free rates ( $R_{f,t}$ ) along with daily Fama-French 5-Factor returns ( $Mkt - RF_t, SMB_t, HML_t, RMW_t, CMA_t$ ) and Carhart Momentum ( $MOM_t$ ) are retrieved directly from the Ken French Data Library (Fama and French, 1993, 2015; Carhart, 1997).

### 3.1.1 Point-in-Time Bloomberg Filing Date Alignment

A common issue in empirical finance studies is lookahead bias caused by indexing corporate accounting metrics to fiscal period end-dates. In practice, financial statements for a fiscal quarter ending on December 31 ( $t_{\text{fiscal}}$ ) are published weeks or months later during the 10-K or 10-Q SEC filing process (López de Prado, 2018).

To eliminate lookahead contamination, my data pipeline (`data_pipeline.py` in `src/`) ingests official Bloomberg publication timestamps (`earnings_announcement_date` and `filing_date`). On any given trading date  $t$ , a quarterly fundamental feature vector  $\mathbf{x}_{i,t}^{\text{fund}}$  is updated to the newly reported accounting state  $Q_\tau$  if and only if:

$$t \geq \text{Announcement Date}(Q_\tau) \quad (3.1)$$

Prior to  $t = \text{Announcement Date}(Q_\tau)$ , the model is strictly limited to using the previously announced fundamental vector  $Q_{\tau-1}$ .

In addition, I enforce a **Staleness Truncation Rule**: a fundamental observation remains active for at most 90 trading days ( $\approx 1$  calendar quarter). If a firm fails to publish a new statement within 120 calendar days ( $\approx 90$  trading days), the forward-fill is truncated and the corresponding fundamental fields are set to missing (NaN), preventing outdated accounting ratios from skewing cross-sectional ranks (Novy-Marx and Velikov, 2016).

### 3.1.2 Survivorship-Bias-Free Dynamic Panel

To guard against survivorship bias, my asset universe tracks historical point-in-time S&P 500 constituent rosters obtained from Bloomberg. Companies removed from the index, acquired, or delisted due to bankruptcy between 2015 and 2024 remain in the daily cross-section up to their final active trading date  $T_{\text{exit}}$  (McLean and Pontiff, 2016; Hou et al., 2020).

Figure 3.1 shows my point-in-time data alignment pipeline and staleness truncation flowchart.



Figure 3.1: Bloomberg Point-in-Time Data Pipeline and Staleness Truncation Flowchart

## 3.2 Feature Space Taxonomy and Selection

My empirical panel relies on **59 variables in total**, split between predictive model inputs and benchmark risk factors:

- **53 Predictive Machine Learning Inputs:** Fed directly into the machine learning and neural network architectures across four modalities: Technical (16), Fundamental (30), Sentiment (2), and Macro (5) (Gu et al., 2020; Cong et al., 2021).
- **6 Benchmark Risk Factor Betas:** Retained for baseline linear regressions and econometric risk adjustments, including the 5 Fama-French systematic factor betas ( $Mkt - RF$ ,  $SMB$ ,  $HML$ ,  $RMW$ ,  $CMA$ ) and Momentum ( $MOM$ ) (Fama and French, 2015; Carhart, 1997).

### 3.2.1 Economic Rationale for Feature Set Size (59 Variables vs. 120+ Zoo Features)

When designing the characteristic space, a key decision is whether to include hundreds of candidate anomaly signals or focus on a curated panel. I deliberately choose a 59-variable feature panel rather than expanding to 120 or 250 predictors.

Including vast numbers of firm characteristics in financial panels introduces practical and statistical problems. Many raw accounting metrics simply restate similar economic signals, such as minor variations of leverage ratios or momentum lookbacks, adding collinearity and noise rather than genuine predictive content (Freyberger et al., 2020; Kozak et al., 2020). In addition, as Harvey et al. (2016) observe, testing hundreds of features increases false-discovery rates. Restricting the predictor set to 53 ML inputs grounded in four economic domains (Technical, Fundamental, Sentiment, and Macro) ensures each variable represents an established economic mechanism (Pástor and Stambaugh, 2003; Novy-Marx, 2013; Cong et al., 2021; Fama and French, 2015). Finally, requiring hundreds of characteristics reduces cross-sectional sample coverage, as

missing fundamental items force the exclusion of valid stock-day observations (Avramov et al., 2023).

Figure 3.2 displays the feature group modality classification.



Figure 3.2: Variable Taxonomy and Feature Group Modality Split (53 ML Inputs + 6 Factor Betas = 59 Total)

### 3.3 Feature Normalization and Selection Protocol

To maintain uniform scaling and stationarity, raw characteristic variables are converted each day into **Cross-Sectional Ranks** mapped to the unit interval  $[0, 1]$  (Freyberger et al., 2020; Gu et al., 2020):

$$r_{i,t,j} = \frac{\text{Rank}(X_{i,t,j}) - 1}{N_t - 1} \in [0, 1] \quad (3.2)$$

Rank-scaling protects models from extreme financial outliers, standardizes variables across different accounting units, and keeps distributions comparable across changing market environments.

Importantly, cross-sectional rank normalization in Equation (3.2) is applied strictly to firm-level characteristics (Technical and Fundamental features). Macroeconomic state variables ( $x_{\text{macro}}$ ) and news/social sentiment indicators are **not** rank-normalized cross-sectionally across stocks on day  $t$  — doing so would reduce their cross-sectional variance to zero, as macro variables are uniform across all assets on a given trading day. Instead, macro state variables are normalized temporally across time or processed directly by the 3-Way Macro Dynamic Gating network.

To guarantee zero lookahead leakage and zero overnight gap contamination, return targets  $y_{i,t+1}$  in training and evaluation are strictly computed from  $P_{i,t+1,\text{open}}$  to  $P_{i,t+2,\text{open}}$  (or  $P_{i,t+1,\text{close}}$  to  $P_{i,t+2,\text{close}}$  following execution at  $t + 1$ ), matching the 1-day execution buffer delay ( $\text{shift}(-2)$ ).

To avoid data leakage, I perform feature selection (`select_features.py` in `scripts/`) using a **three-stage protocol applied exclusively to pre-2020 training data** (146,959 observations) (López de Prado, 2018). The first stage removes features with missing rates above 10%. The second stage prunes redundant variables exhibiting pairwise Spearman correlations  $|\rho_{j,k}| > 0.85$ . In the third stage, an in-sample Random Forest regressor (Breiman, 2001) ranks remaining candidate variables, selecting the top 53 predictive features.

Figure 3.3 presents the correlation matrix of selected features, and Figure 3.4 shows the Random Forest importance rankings.




Figure 3.3: Spearman Rank Correlation Heatmap of Selected 53 ML Features

Figure 3.4 reports the in-sample permutation feature importances calculated strictly on pre-2020 training data to prevent out-of-sample lookahead leakage. A classic methodological issue in tree-based feature selection is that raw Gini impurity measures over-attribute importance to continuous, un-normalized variables (Strobl et al., 2007). Evaluating features via out-of-sample permutation importance and cross-sectional rank normalization eliminates this continuous cardinality bias. Technical momentum indicators ( $\text{vol}_{126d}$ ,  $\text{mom}_{21d}$ ,  $\text{rsi}_{14d}$ ) and fundamental metrics ( $\text{pe}_{\text{ratio}}$ ,  $\text{roi}$ ,  $\text{quality}$ ) emerge as




Figure 3.4: Stage 3 In-Sample Permutation Feature Importance Rankings across Modalities

the dominant predictive drivers, while sentiment indicators provide moderate secondary refinement, aligning fully with downstream SHAP values (Figure 4.2) and modality ablation evidence (Table 4.3).

### 3.4 Baseline Econometric & Machine Learning Formulations

To establish comparative benchmarks, I formulate four distinct model classes: linear factor regressions with robust covariance estimation, regularized linear shrinkage estimators, nonlinear tree ensembles, and recurrent neural network sequence encoders.

### 3.4.1 Fama-MacBeth 2-Pass Linear Regression & Newey-West HAC Estimation

The benchmark linear asset pricing model follows the two-pass cross-sectional regression procedure of Fama and MacBeth (1973), incorporating 1-day lagged factor loadings to prevent lookahead bias and heteroskedasticity and autocorrelation consistent (HAC) covariance matrices (Newey and West, 1987).

#### Pass 1: Rolling Time-Series Factor Loading Estimation

For each individual stock  $i \in \{1, \dots, N_t\}$  on trading day  $t$ , factor loadings (betas) are estimated over a rolling historical lookback window of  $W = 252$  trading days (one calendar year). Let  $\mathbf{R}_{i,t} = [R_{i,t-W+1}, R_{i,t-W+2}, \dots, R_{i,t}]^T \in \mathbb{R}^{W \times 1}$  denote the vector of historical daily excess returns for asset  $i$ , and let  $\mathbf{F}_t \in \mathbb{R}^{W \times (K+1)}$  represent the matrix of  $K$  benchmark factor returns plus a column of ones for the intercept:

$$\mathbf{F}_t = \begin{bmatrix} 1 & F_{1,t-W+1} & F_{2,t-W+1} & \cdots & F_{K,t-W+1} \\ 1 & F_{1,t-W+2} & F_{2,t-W+2} & \cdots & F_{K,t-W+2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & F_{1,t} & F_{2,t} & \cdots & F_{K,t} \end{bmatrix} \quad (3.3)$$

The first-pass time series Ordinary Least Squares (OLS) regression for stock  $i$  is formulated as:

$$\mathbf{R}_{i,t} = \mathbf{F}_t \boldsymbol{\beta}_{i,t} + \boldsymbol{\epsilon}_{i,t} \quad (3.4)$$

where  $\boldsymbol{\beta}_{i,t} = [\alpha_{i,t}, \beta_{i,1,t}, \beta_{i,2,t}, \dots, \beta_{i,K,t}]^T \in \mathbb{R}^{(K+1) \times 1}$ . The OLS point estimator is obtained by minimizing the sum of squared residuals:

$$\hat{\boldsymbol{\beta}}_{i,t} = \arg \min_{\boldsymbol{\beta}} (\mathbf{R}_{i,t} - \mathbf{F}_t \boldsymbol{\beta})^T (\mathbf{R}_{i,t} - \mathbf{F}_t \boldsymbol{\beta}) = (\mathbf{F}_t^T \mathbf{F}_t)^{-1} \mathbf{F}_t^T \mathbf{R}_{i,t} \quad (3.5)$$

To enforce zero-lookahead protection, factor loadings  $\hat{\boldsymbol{\beta}}_{i,t}$  estimated using data up to day  $t$  are indexed to predict returns at day  $t + 1$ .

## Pass 2: Daily Cross-Sectional Factor Risk Premia Estimation

On each trading day  $t + 1$ , a cross-sectional OLS regression is estimated across all active constituent assets  $N_{t+1}$  using the 1-day lagged factor loadings  $\hat{\mathbf{B}}_t = [\mathbf{1}, \hat{\boldsymbol{\beta}}_{1,t}, \hat{\boldsymbol{\beta}}_{2,t}, \dots, \hat{\boldsymbol{\beta}}_{N_{t+1},t}]^T \in \mathbb{R}^{N_{t+1} \times (K+1)}$ :

$$\mathbf{R}_{t+1} = \hat{\mathbf{B}}_t \boldsymbol{\gamma}_{t+1} + \boldsymbol{\eta}_{t+1} \quad (3.6)$$

where  $\mathbf{R}_{t+1} = [R_{1,t+1}, R_{2,t+1}, \dots, R_{N_{t+1},t+1}]^T \in \mathbb{R}^{N_{t+1} \times 1}$  is the cross-section of realized forward returns,  $\boldsymbol{\gamma}_{t+1} = [\gamma_{0,t+1}, \gamma_{1,t+1}, \dots, \gamma_{K,t+1}]^T \in \mathbb{R}^{(K+1) \times 1}$  represents the vector of daily factor risk premia (and cross-sectional intercept  $\gamma_{0,t+1}$ ), and  $\boldsymbol{\eta}_{t+1}$  is the vector of cross-sectional pricing errors.

The closed-form daily cross-sectional risk premium estimator is given by:

$$\hat{\boldsymbol{\gamma}}_{t+1} = \left( \hat{\mathbf{B}}_t^T \hat{\mathbf{B}}_t \right)^{-1} \hat{\mathbf{B}}_t^T \mathbf{R}_{t+1} \quad (3.7)$$

## Time-Series Aggregation & Step-by-Step Newey-West HAC Proof

Over an evaluation period of  $T$  trading days, the overall factor risk premium vector is estimated as the sample mean across daily cross-sectional estimates:

$$\bar{\boldsymbol{\gamma}} = \frac{1}{T} \sum_{t=1}^T \hat{\boldsymbol{\gamma}}_t \quad (3.8)$$

Because daily factor risk premia estimates  $\hat{\boldsymbol{\gamma}}_t$  suffer from serial correlation (autocorrelation) due to overlapping return measurement horizons and conditional heteroskedasticity in asset volatility, standard i.i.d. variance estimators underestimate true standard errors, inflating test statistics.



To derive a consistent estimator for the asymptotic covariance matrix of  $\bar{\gamma}$ , let  $\mathbf{u}_t = \hat{\gamma}_t - \bar{\gamma} \in \mathbb{R}^{(K+1) \times 1}$  denote the vector of mean-zero risk premium innovations. The variance of the sample average factor risk premium is:

$$\text{Var}(\bar{\gamma}) = \text{Var} \left( \frac{1}{T} \sum_{t=1}^T \hat{\gamma}_t \right) = \frac{1}{T^2} \sum_{t=1}^T \sum_{s=1}^T \mathbb{E} [\mathbf{u}_t \mathbf{u}_s^T] \quad (3.9)$$

Defining the sample autocovariance matrix at lag  $j$  as:

$$\hat{\mathbf{\Gamma}}_j = \frac{1}{T} \sum_{t=j+1}^T \mathbf{u}_t \mathbf{u}_{t-j}^T, \quad \text{for } j \geq 0 \quad (3.10)$$

Equation (3.9) can be expanded into the sum of zero-lag covariance and off-diagonal cross-lag covariances:

$$\text{Var}(\bar{\gamma}) = \frac{1}{T} \left[ \hat{\mathbf{\Gamma}}_0 + \sum_{j=1}^{T-1} \left( \hat{\mathbf{\Gamma}}_j + \hat{\mathbf{\Gamma}}_j^T \right) \right] \quad (3.11)$$

Unweighted truncation of lag covariances at finite lag  $L < T$  can lead to non-positive semi-definite covariance matrices in finite samples. The Newey and West (1987) HAC covariance matrix estimator guarantees positive semi-definiteness by weighting lag autocovariances with the triangular Bartlett kernel  $w(j, L)$ :

$$\hat{\mathbf{\Omega}}_{\text{HAC}} = \hat{\mathbf{\Gamma}}_0 + \sum_{j=1}^L w(j, L) \left( \hat{\mathbf{\Gamma}}_j + \hat{\mathbf{\Gamma}}_j^T \right) \quad (3.12)$$

where the Bartlett kernel weights are defined as:

$$w(j, L) = 1 - \frac{j}{L+1}, \quad j = 1, 2, \dots, L \quad (3.13)$$

and the lag truncation bandwidth  $L$  is set according to the bandwidth selection rule  $L = \lfloor 4(T/100)^{2/9} \rfloor$ .

*Proof of Guaranteed Positive Semi-Definiteness.* To prove that  $\hat{\mathbf{\Omega}}_{\text{HAC}} \succeq 0$ , express the Bartlett kernel as the spectral window resulting from the convolution of a rectangular

sequence with itself. Define  $\mathbf{u}_t = \mathbf{0}$  for  $t \notin \{1, \dots, T\}$ . Consider the smoothed sequence  $\mathbf{s}_t = \frac{1}{\sqrt{L+1}} \sum_{m=0}^L \mathbf{u}_{t+m}$ . The sample covariance of  $\mathbf{s}_t$  is positive semi-definite:

$$\mathbf{M} = \frac{1}{T} \sum_{t=-\infty}^{\infty} \mathbf{s}_t \mathbf{s}_t^T = \frac{1}{T(L+1)} \sum_{t=-\infty}^{\infty} \left( \sum_{m=0}^L \mathbf{u}_{t+m} \right) \left( \sum_{n=0}^L \mathbf{u}_{t+n}^T \right) \geq 0 \quad (3.14)$$

Expanding the double summation yields:

$$\begin{aligned} \mathbf{M} &= \frac{1}{T(L+1)} \sum_{m=0}^L \sum_{n=0}^L \sum_{t=-\infty}^{\infty} \mathbf{u}_{t+m} \mathbf{u}_{t+n}^T = \frac{1}{T(L+1)} \sum_{m=0}^L \sum_{n=0}^L T \hat{\mathbf{\Gamma}}_{m-n} \\ &= \frac{1}{L+1} \sum_{k=-L}^L (L+1-|k|) \hat{\mathbf{\Gamma}}_k = \hat{\mathbf{\Gamma}}_0 + \sum_{j=1}^L \left( 1 - \frac{j}{L+1} \right) (\hat{\mathbf{\Gamma}}_j + \hat{\mathbf{\Gamma}}_j^T) = \hat{\mathbf{\Omega}}_{\text{HAC}} \end{aligned} \quad (3.15)$$

$$(3.16)$$

Since  $\mathbf{M}$  is constructed as a sum of outer products  $\mathbf{s}_t \mathbf{s}_t^T$ , for any vector  $\mathbf{a} \in \mathbb{R}^{K+1}$ ,  $\mathbf{a}^T \hat{\mathbf{\Omega}}_{\text{HAC}} \mathbf{a} = \frac{1}{T} \sum_t (\mathbf{a}^T \mathbf{s}_t)^2 \geq 0$ . Thus,  $\hat{\mathbf{\Omega}}_{\text{HAC}}$  is positive semi-definite.  $\square$

The Newey-West adjusted standard error for the  $k$ -th factor risk premium estimate  $\bar{\gamma}_k$  is:

$$\text{SE}_{\text{NW}}(\bar{\gamma}_k) = \sqrt{\frac{1}{T} [\hat{\mathbf{\Omega}}_{\text{HAC}}]_{kk}} \quad (3.17)$$

and the test statistic for factor significance testing is evaluated as  $t(\bar{\gamma}_k) = \frac{\bar{\gamma}_k}{\text{SE}_{\text{NW}}(\bar{\gamma}_k)}$ .